

Computer Vision

Lecture 12 – Diverse Topics in Computer Vision

Prof. Dr.-Ing. Andreas Geiger

Autonomous Vision Group

University of Tübingen / MPI-IS



e l l i s
European Laboratory for Learning and Intelligent Systems

In this class, we have learned a lot ..

- ▶ History of computer vision
- ▶ Primitives and transformations
- ▶ Geometric image formation
- ▶ Photometric image formation
- ▶ Camera calibration
- ▶ SIFT detection and matching
- ▶ Epipolar geometry
- ▶ Structure-from-Motion
- ▶ Bundle adjustment
- ▶ Image rectification
- ▶ Block matching
- ▶ Siamese networks
- ▶ Probabilistic graphical models
- ▶ Belief propagation
- ▶ Structured prediction
- ▶ Deep structured models
- ▶ Volumetric 3D reconstruction
- ▶ Optical flow
- ▶ Rendering equation
- ▶ Shape-from-Shading
- ▶ Photometric stereo
- ▶ Multi-view stereo
- ▶ Structured light
- ▶ Monocular depth
- ▶ Volumetric fusion
- ▶ Marching cubes
- ▶ Implicit neural representations
- ▶ Occupancy networks
- ▶ Differentiable rendering
- ▶ Novel view synthesis
- ▶ Neural radiance fields
- ▶ Generative radiance fields
- ▶ Image classification
- ▶ Semantic segmentation
- ▶ Object detection
- ▶ Average precision
- ▶ Instance segmentation
- ▶ Data annotation
- ▶ Self-supervised learning
- ▶ Self-supervised depth/flow
- ▶ Pretext tasks
- ▶ Contrastive learning

.. but there is a lot more:

- ▶ 3D feature learning
- ▶ Shape abstraction
- ▶ Neural rendering
- ▶ Image correspondence
- ▶ SLAM, VO & localization
- ▶ Morphable 3D models
- ▶ 3D detection
- ▶ Action recognition
- ▶ Adversarial attacks
- ▶ Face recognition
- ▶ Gaze estimation
- ▶ Human body models
- ▶ Event cameras
- ▶ Omnidirectional vision
- ▶ Deblurring, denoising, dehazing
- ▶ Superresolution
- ▶ Lightfields
- ▶ Neuromorphic sensors
- ▶ Datasets and benchmarks
- ▶ Network pruning
- ▶ Few-shot learning
- ▶ Open set recognition
- ▶ Image/video retrieval
- ▶ Image synthesis, GANs, VAEs
- ▶ Neural style transfer
- ▶ Explainable AI
- ▶ Ethics, privacy & fairness
- ▶ Deepfakes
- ▶ Video prediction
- ▶ Saliency prediction
- ▶ Bio/medical imaging
- ▶ Object tracking
- ▶ 6D Pose estimation
- ▶ Active learning
- ▶ Disentangled representations
- ▶ Scene graphs
- ▶ Domain adaptation
- ▶ Geometric deep learning
- ▶ Ego-centric vision
- ▶ Vision and language
- ▶ Robot vision
- ▶ Surveillance

Credits

- ▶ **Justin Johnson** – Convolutional Neural Networks + Neural Style Transfer
http://web.stanford.edu/class/cs20si/lectures/slides_06.pdf
- ▶ **Leon Gatys** – Image Style Transfer Using Convolutional Neural Networks
<https://www.youtube.com/watch?v=UFffxcCQMPQ>
- ▶ **Michael Black** – SMPL made Simple
<https://smpl-made-simple.is.tue.mpg.de/>
<https://www.youtube.com/watch?v=rzpiSYTrRU0>
- ▶ **Matthias Niessner** – Deepfakes Creation and Detection
<https://www.youtube.com/watch?v=-Xv2IRs2-KA>

Agenda

12.1 Input Optimization

12.2 Compositional Models

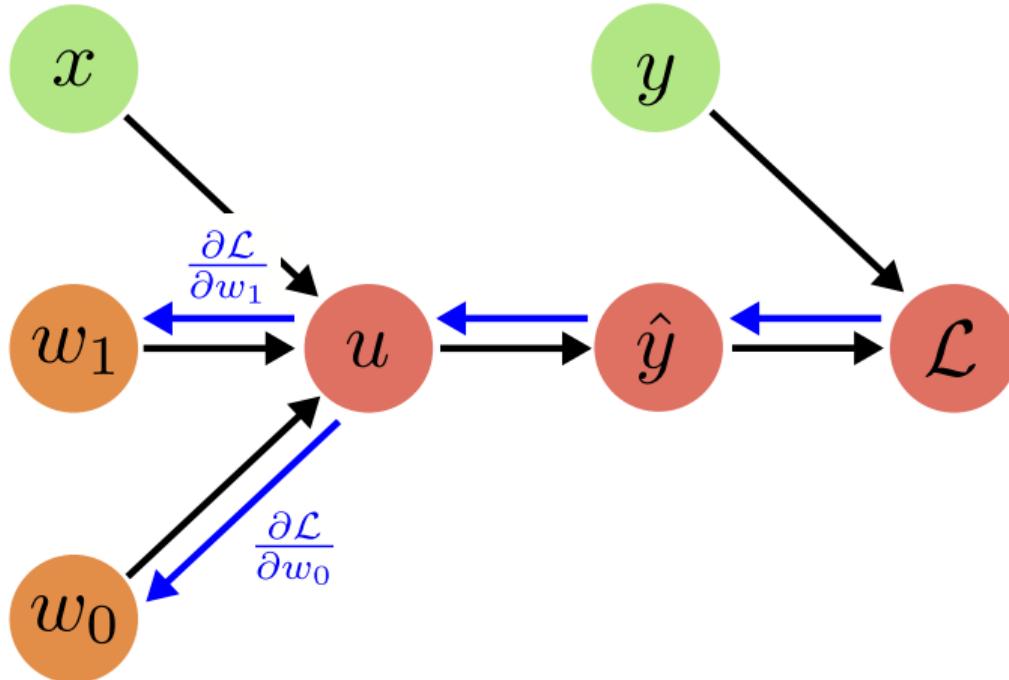
12.3 Human Body Models

12.4 Deepfakes

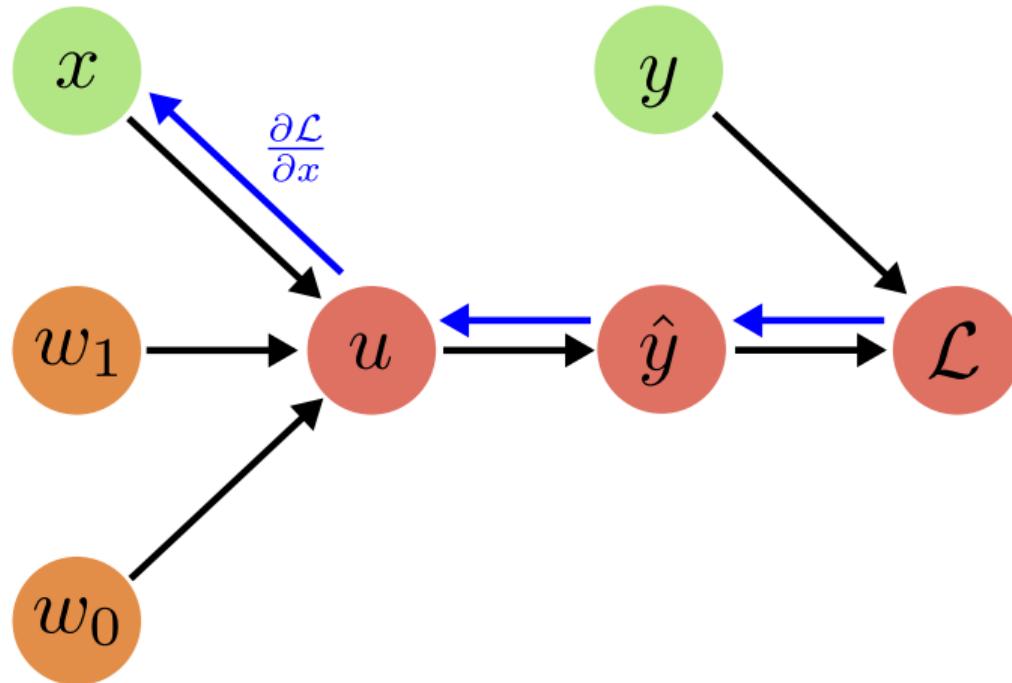
12.1

Input Optimization

Weight Optimization vs. Input Optimization



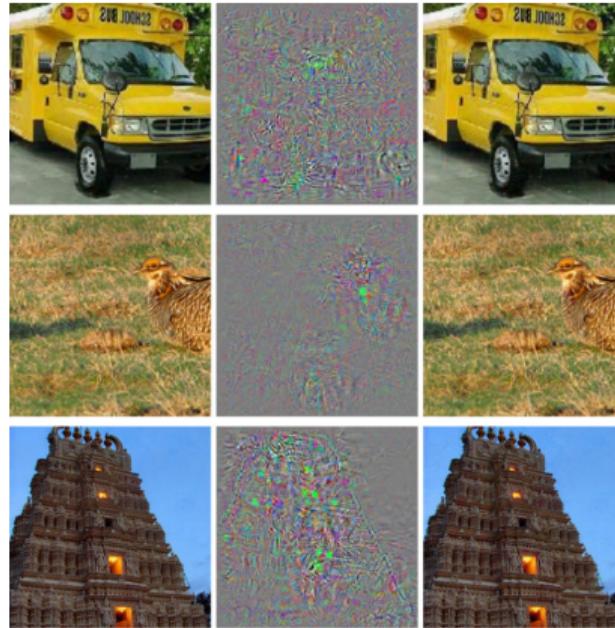
Weight Optimization vs. Input Optimization



- We will see two examples in this unit: **Adversarial attacks** and **style transfer**

Adversarial Attacks

Adversarial Attacks on Image Classification



L-BFGS Attack:

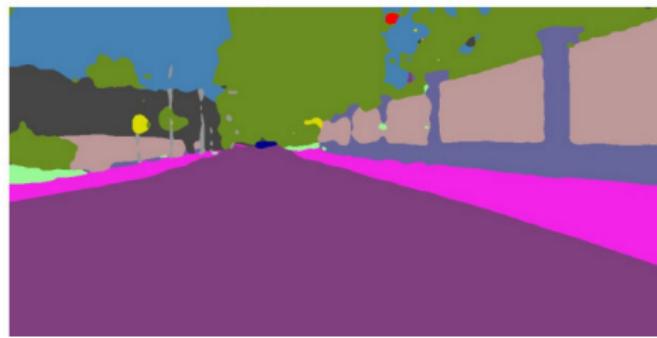
- ▶ Given classifier $f : \mathbb{R}^m \rightarrow \{1, \dots, L\}$
- ▶ Find **adversarial example** for image x

$$x + \operatorname{argmin}_{\Delta x} \{\|\Delta x\|_2 : f(x + \Delta x) = y_t\}$$

where y_t denotes the target class

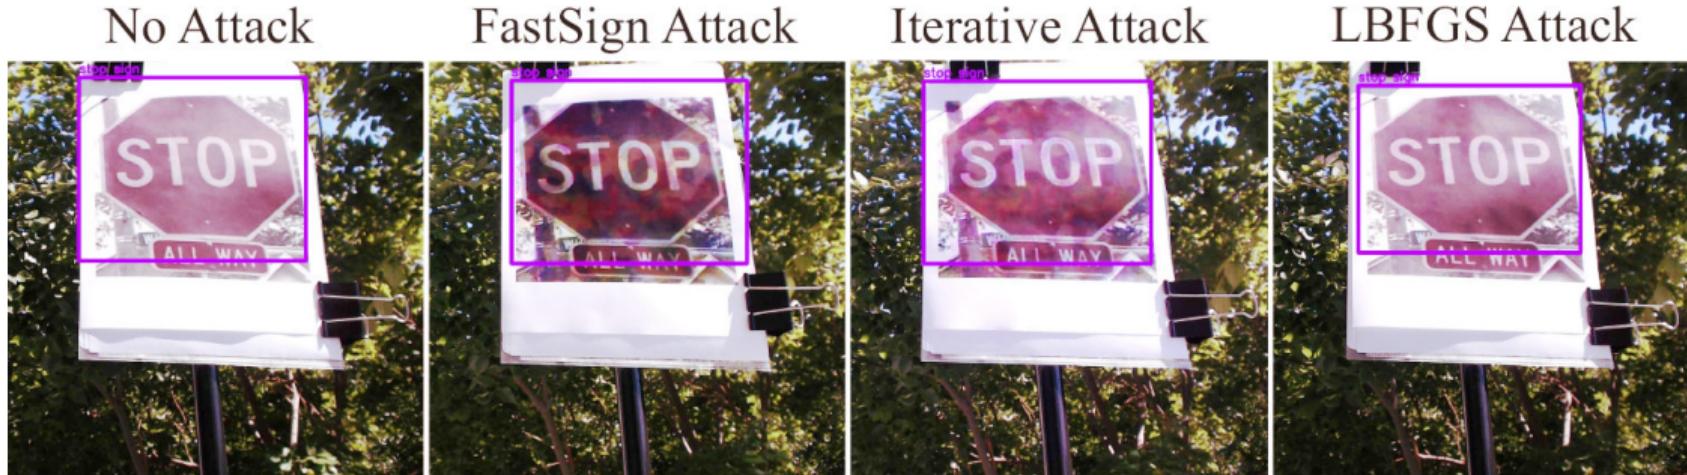
- ▶ All images in the right column classified as “ostrich”, manipulations (Δx) are magnified

Adversarial Attacks on Semantic Segmentation



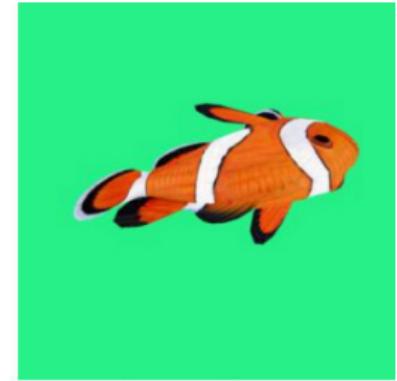
- ▶ Attack on **semantic segmentation** manipulates label map

Physical Adversarial Attacks



"Adversarial perturbation methods applied to **stop sign detection** only work in carefully chosen situations, and our preliminary experiment shows that we might **not need to worry** about it in many real circumstances, specifically with autonomous vehicles."

Robust Adversarial Attacks



- ▶ Demonstrate existence of **robust adversarial examples** in the physical world
- ▶ Maximize **expectation over transformation** \mathcal{T} (EOT):

$$\underset{x'}{\operatorname{argmax}} \mathbb{E}_{t \sim \mathcal{T}} [\log P(y_t | t(x')) - \lambda \| (t(x') - t(x)) \|_2]$$

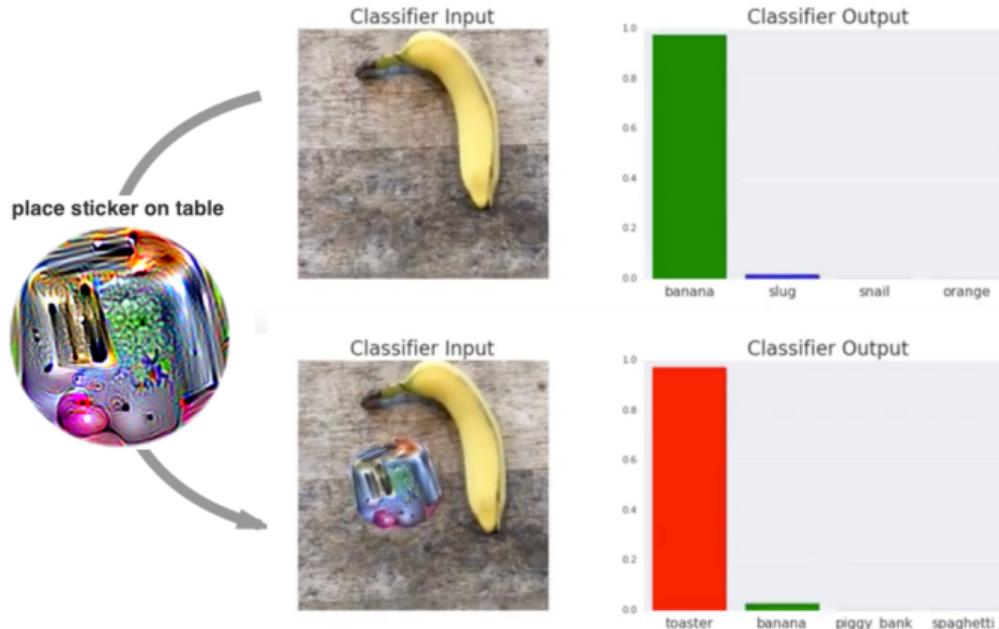
- ▶ Larger distributions require larger perturbations

Robust Adversarial Attacks



- ▶ **Robust adversarial example** designed to mimic “graffiti”

Adversarial Patch Attacks



- ▶ **Patch attacks** use EOT idea, but additionally also optimize across **many images**
- ▶ Easy to apply in real-world settings (attaching patch to an object)

Attacking Optical Flow



Given an image pair from a video sequence ...

Attacking Optical Flow



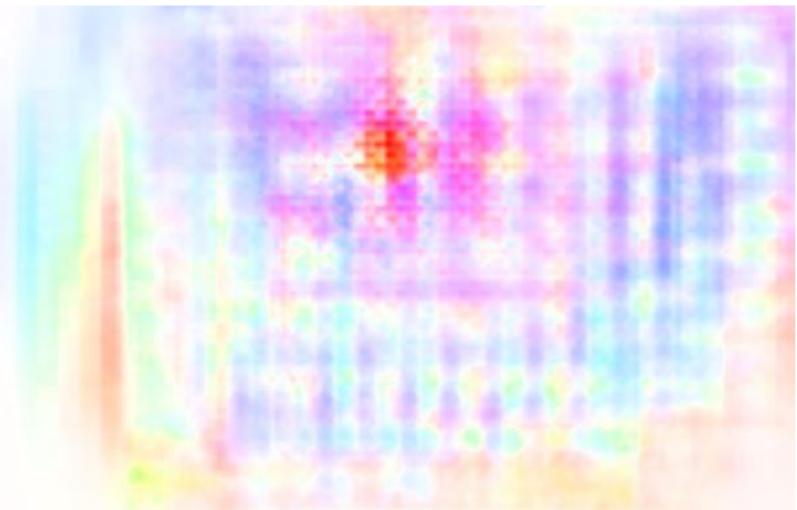
... FlowNet2 predicts a smooth optical flow field.

Attacking Optical Flow



Can we obtain a small ($< 1\%$ area) attack patch \hat{p} ...

Attacking Optical Flow



... that successfully attacks the optical flow network?

Attacking Optical Flow

Let $F(I, I')$ denote an **optical flow network** and \mathcal{I} a dataset of frame pairs $\{(I, I')\}$.

Let $A(I, p, t, l)$ denote **image** I with **patch** p transformed by t inserted at location l .

Let \mathcal{T} denote a distribution over affine 2D transformations.

Let \mathcal{L} denote a (uniform) distribution over the image domain.

Goal: Find a patch \hat{p} such that

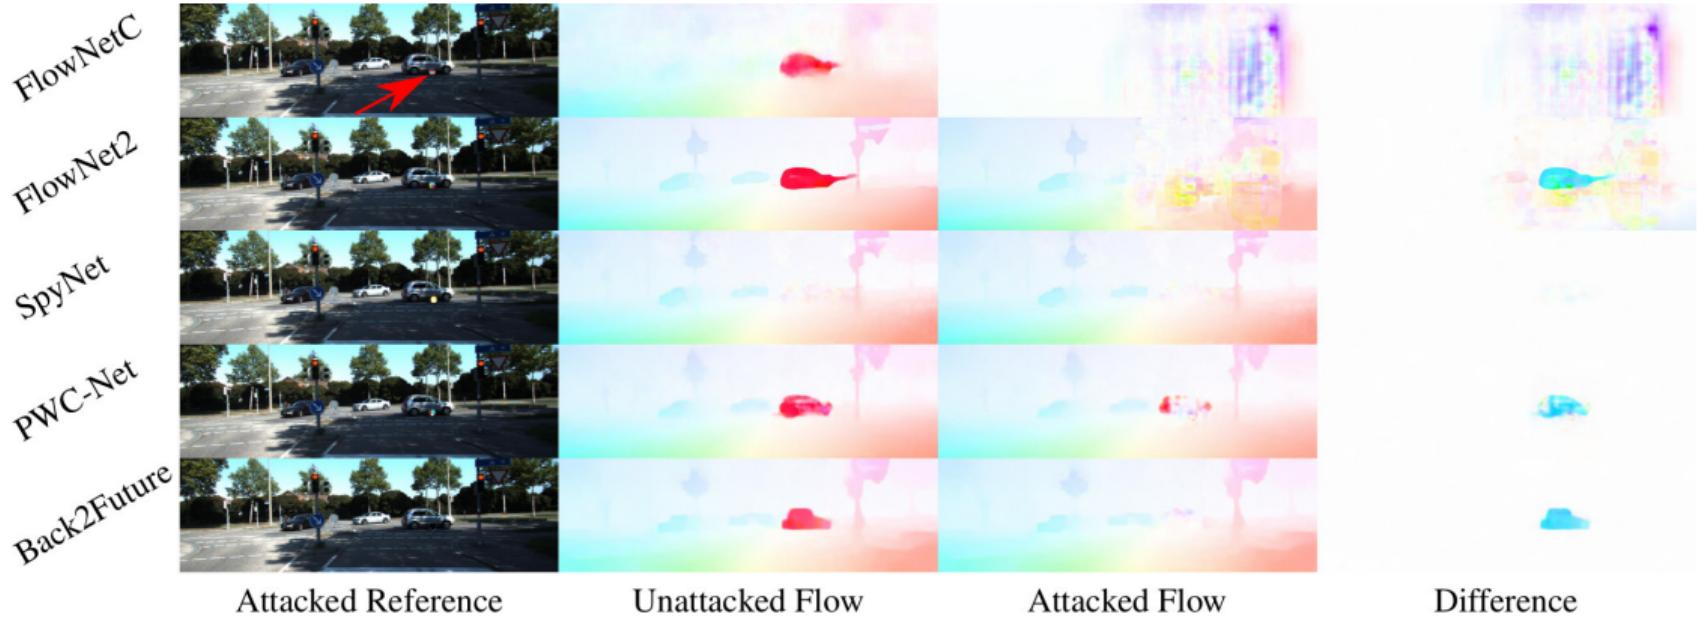
$$\hat{p} = \operatorname{argmin}_p \mathbb{E}_{(I, I') \sim \mathcal{I}, t \sim \mathcal{T}, l \sim \mathcal{L}} \left[\frac{(u, v) \cdot (\tilde{u}, \tilde{v})}{\|(u, v)\|_2 \cdot \|(\tilde{u}, \tilde{v})\|_2} \right]$$

$$\text{with } (u, v) = F(I, I')$$

$$(\tilde{u}, \tilde{v}) = F(A(I, p, t, l), A(I', p, t, l))$$

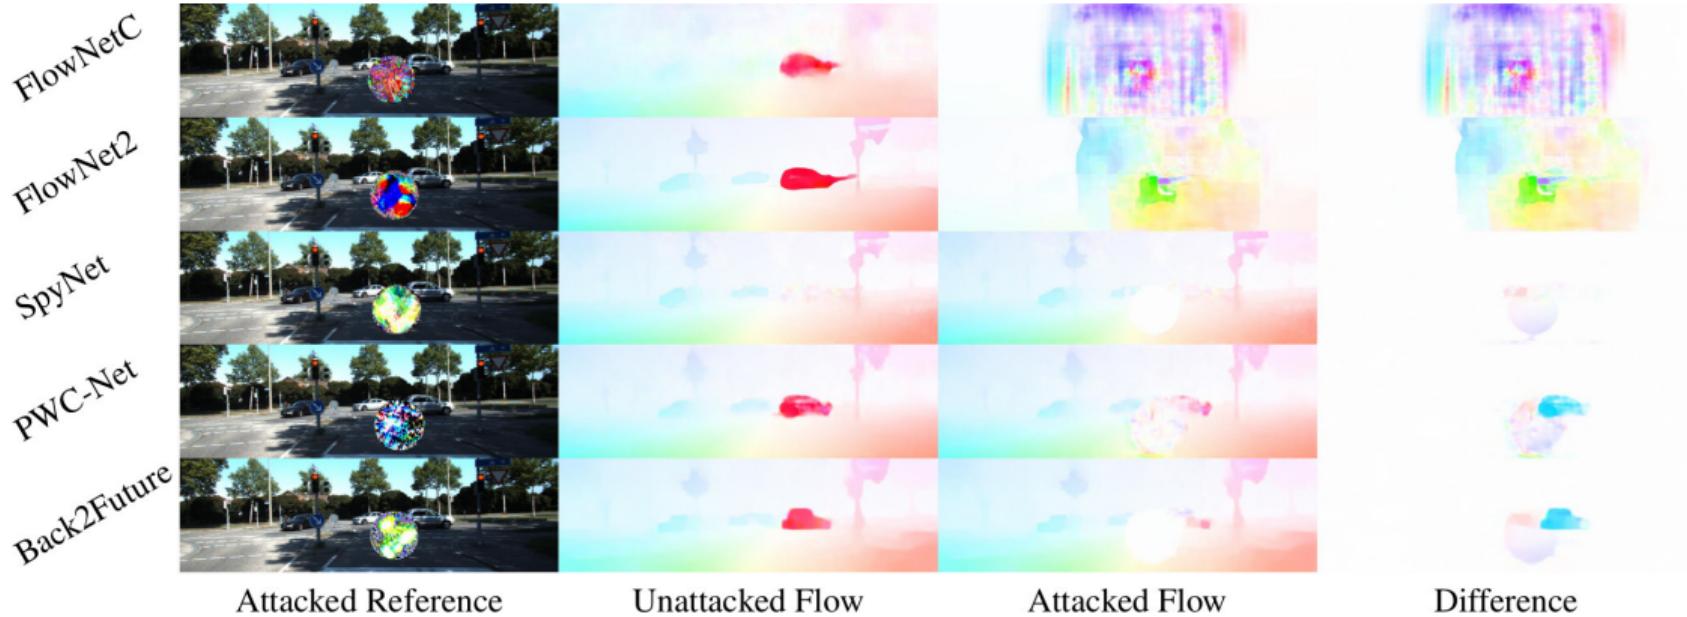
Intuition: Find patch \hat{p} that reverses the direction of the optical flow.

White-Box Attacks



- Attacks **extend beyond region** of patch for encoder-decoder architectures

White-Box Attacks

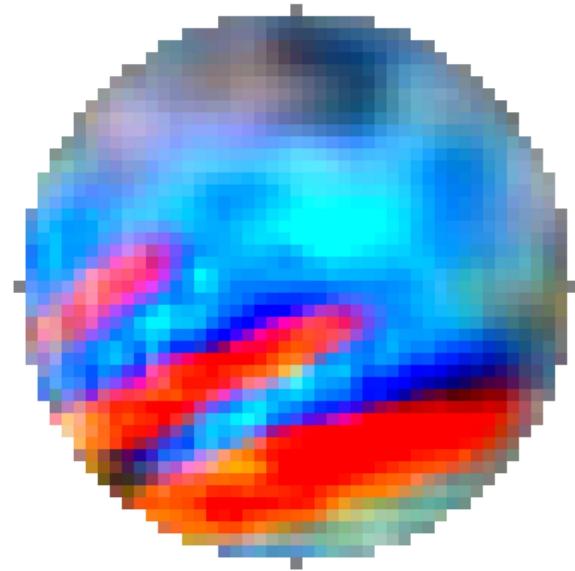


- ▶ Attacks **extend beyond region** of patch for encoder-decoder architectures

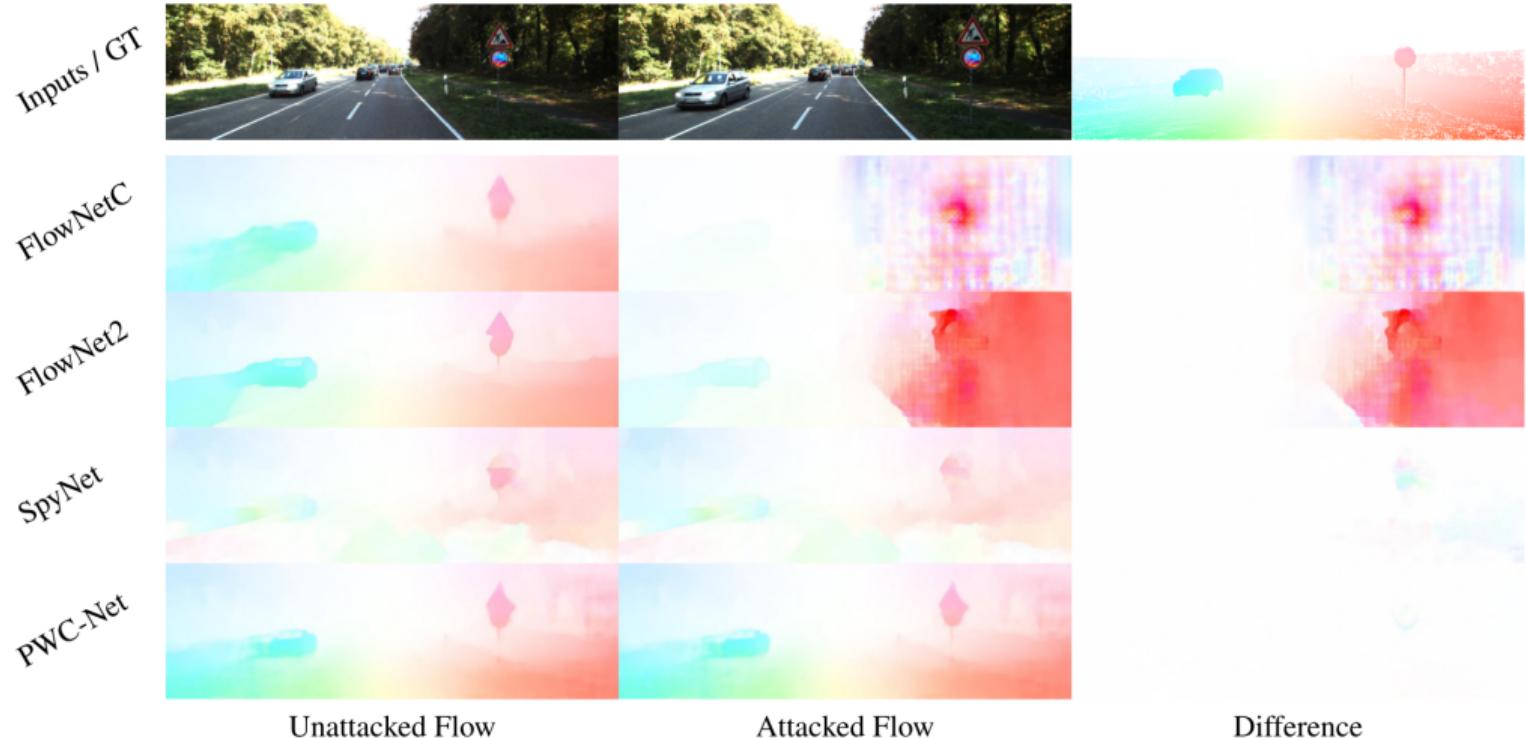
Black-Box Attacks

Setting:

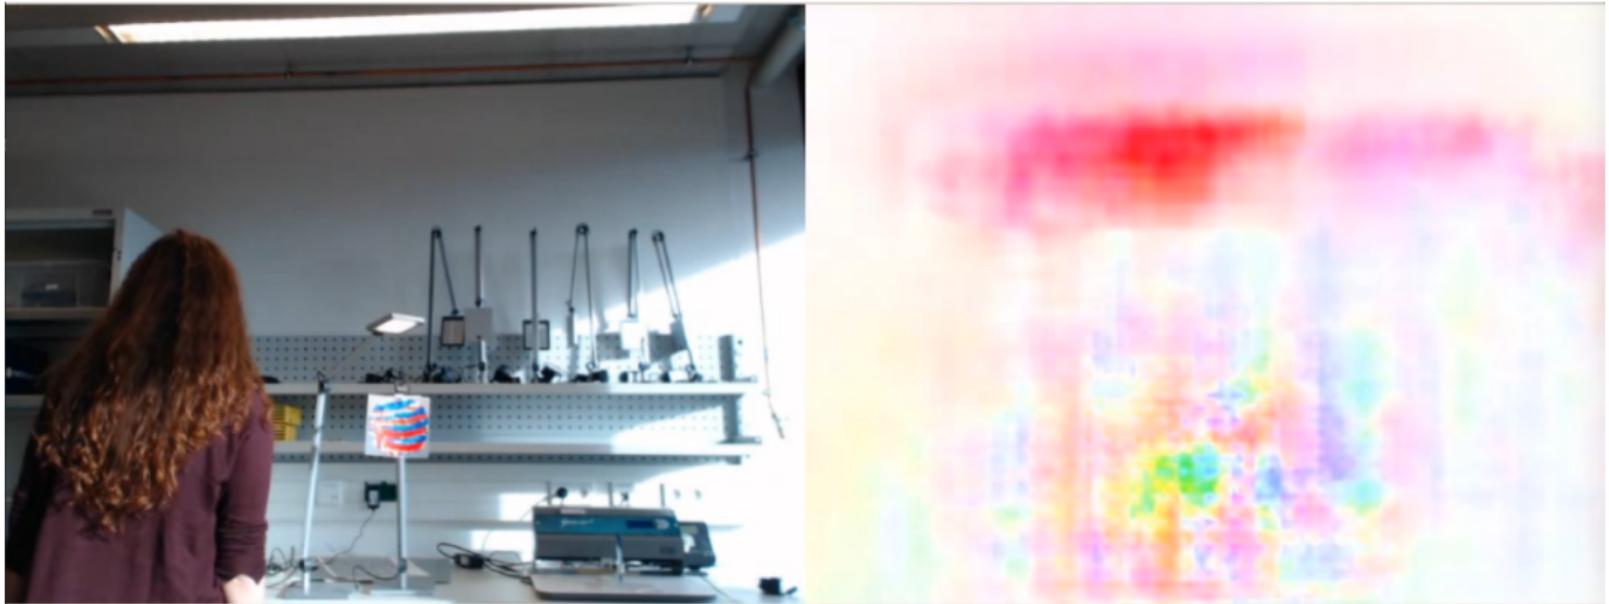
- ▶ Patch optimized for two networks
(FlowNet2 and PWCNet)
- ▶ Used to attack unseen network



Black-Box Attacks



Real-World Attack



- ▶ Printed patch attached to desk lamp

Defenses against Adversarial Attacks

An entire field evolved around **defending adversarial attacks**:

- ▶ **Gradient Masking/Obfuscation:** Since most attack algorithms are based on gradient information of the classifier, masking or hiding the gradients will confound the adversaries. [Papernot, 2016] [Song, 2017] [Buckmann, 2018]
- ▶ **Robust Optimization:** Re-learning a DNN classifier's parameters can increase its robustness. The trained classifier will correctly classify the subsequently generated adversarial examples. [Goodfellow, 2014] [Madry, 2017] [Hein, 2017]
- ▶ **Adversarial Example Detection:** Study the distribution of natural/benign examples, detect adversarial examples and disallow them as input to the classifier. [Carlini, 2017] [Grosse, 2017] [Metzen, 2017]

Neural Style Transfer

Neural Style Transfer



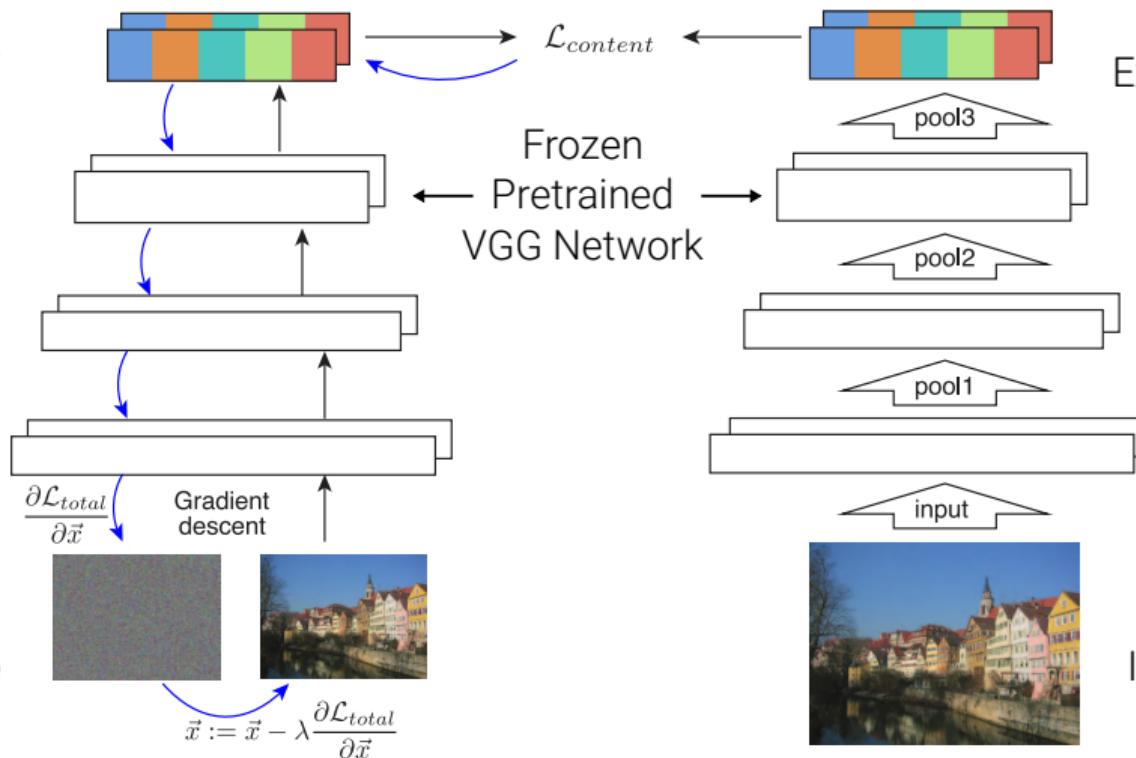
Neural Style Transfer:

- ▶ Synthesize an image with the **content** of one image and the **style** of another
- ▶ Start with **random image**, optimize content and style loss: $\mathcal{L} = \mathcal{L}_{\text{content}} + \mathcal{L}_{\text{style}}$
- ▶ Exploits **VGG** network pre-trained on **ImageNet** as feature extractor

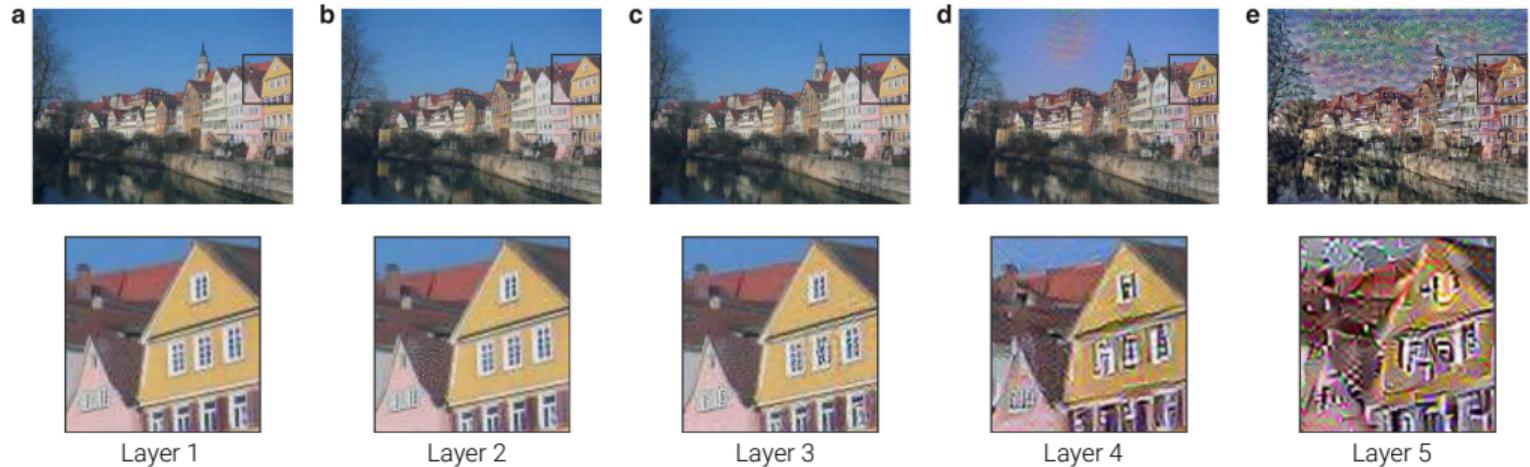
Content Reconstruction

Step 4:

Optimize Image
to minimize
Content Loss



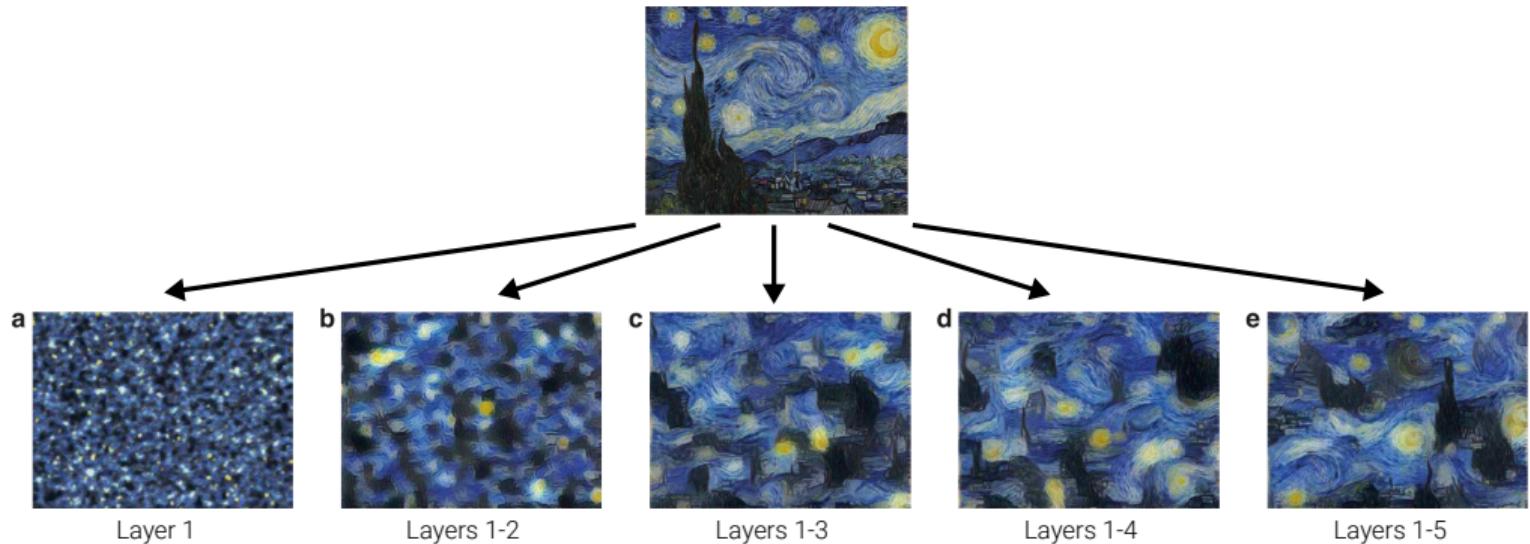
Content Reconstruction



$$\mathcal{L}_{\text{content}}(l) = \|\mathbf{F}_l - \mathbf{P}_l\|_2^2$$

- ▶ \mathbf{F}_l : Features of **generated** image at VGG layer l
- ▶ \mathbf{P}_l : Features of **content** image at VGG layer l

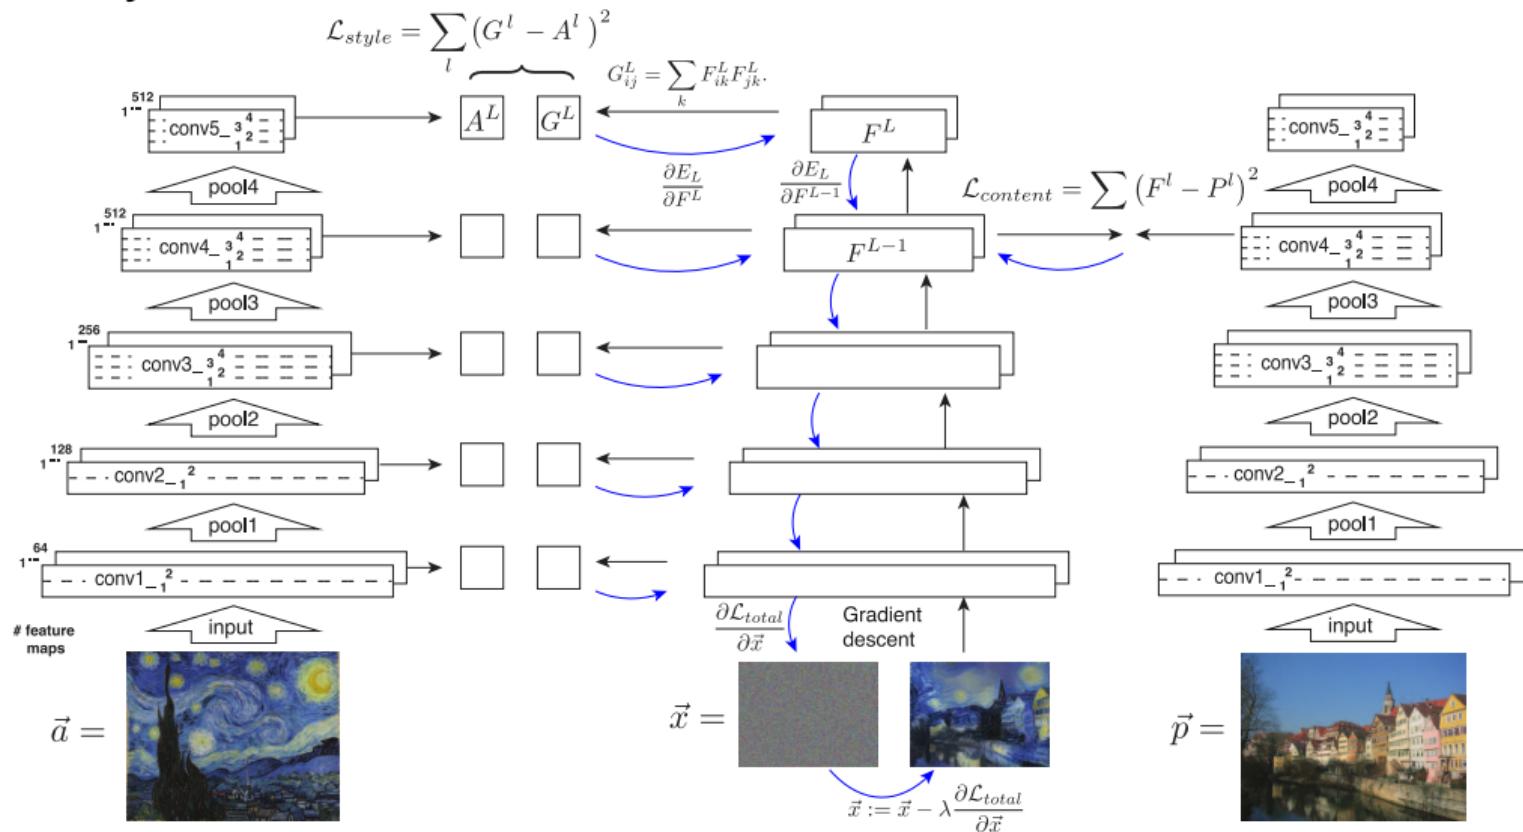
Style Reconstruction



$$\mathcal{L}_{\text{style}}(l) = \|\mathbf{G}_l - \mathbf{A}_l\|_2^2 \quad G_{ij}^l = \sum_{\mathbf{p} \in \Omega} F_i^l(\mathbf{p}) F_j^l(\mathbf{p})$$

- $\mathbf{G}_l, \mathbf{A}_l$: **Gram matrices** (feature correlations) of generated and style image

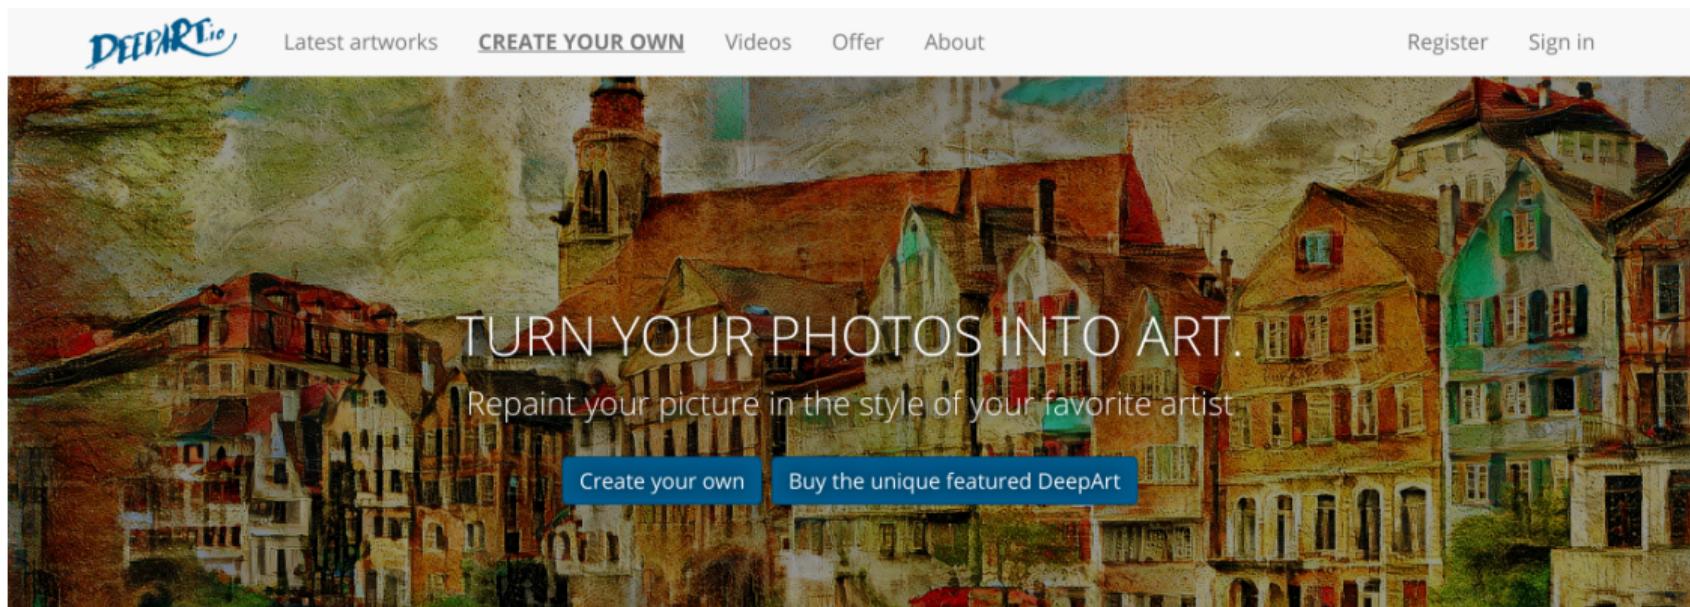
Neural Style Transfer



Neural Style Transfer



Neural Style Transfer

The banner features a painting of a European town with colorful buildings and a church tower. Overlaid text reads: "TURN YOUR PHOTOS INTO ART." and "Repaint your picture in the style of your favorite artist." Below the text are two blue buttons: "Create your own" and "Buy the unique featured DeepArt".

DEEPART.io

Latest artworks [CREATE YOUR OWN](#) Videos Offer About

Register Sign in

TURN YOUR PHOTOS INTO ART.

Repaint your picture in the style of your favorite artist

Create your own Buy the unique featured DeepArt

<https://deepart.io/>

12.2

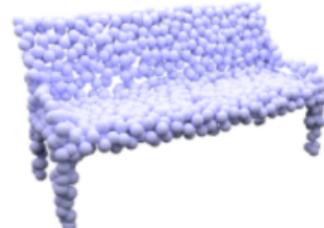
Compositional Models

Shape Abstraction

3D Representations



Voxels



Points



Meshes

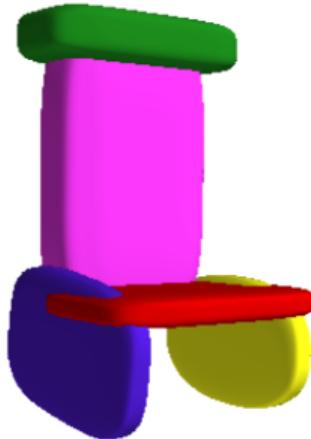
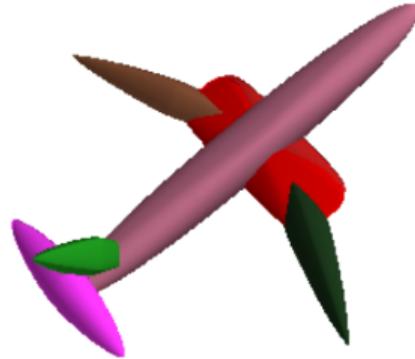


Implicit

Limitations of existing 3D representations:

- ▶ Large number of unstructured geometric elements
- ▶ Do not convey **semantic information** (parts, functionality, etc.)

3D Primitives



Primitive-based 3D Representations:

- ▶ **Few primitives** required to represent a 3D object
- ▶ Convey **semantic information** (parts, functionality, etc.)
- ▶ Challenges: Variable number of primitives, few annotated datasets

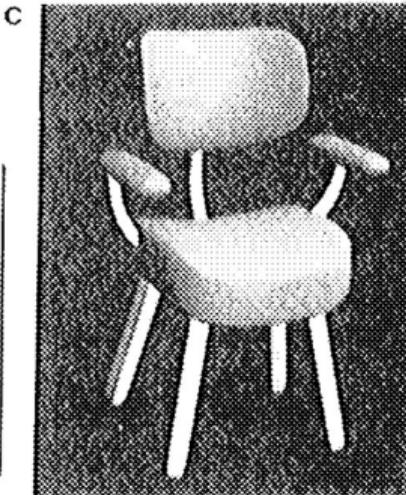
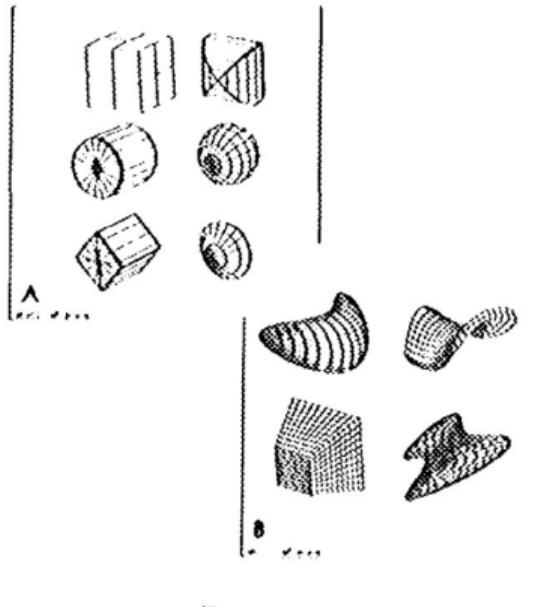
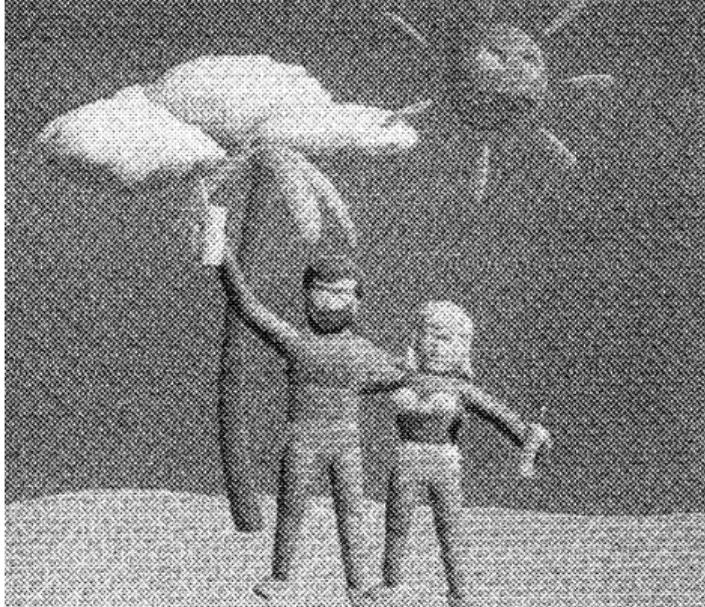
3D Shape Abstraction

Goal:

- ▶ Learn 3D shape abstraction
- ▶ Infer number of primitives
- ▶ No supervision at primitive level
- ▶ Input: point cloud or image
- ▶ Supervision: point cloud (no labels)

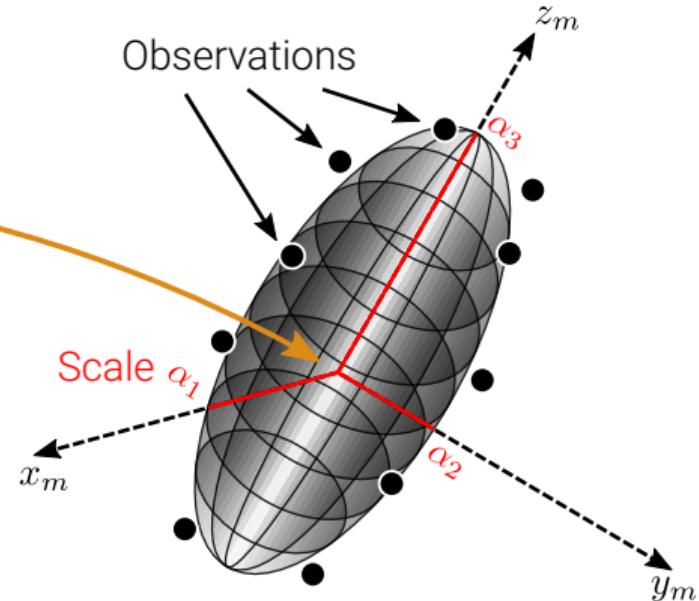
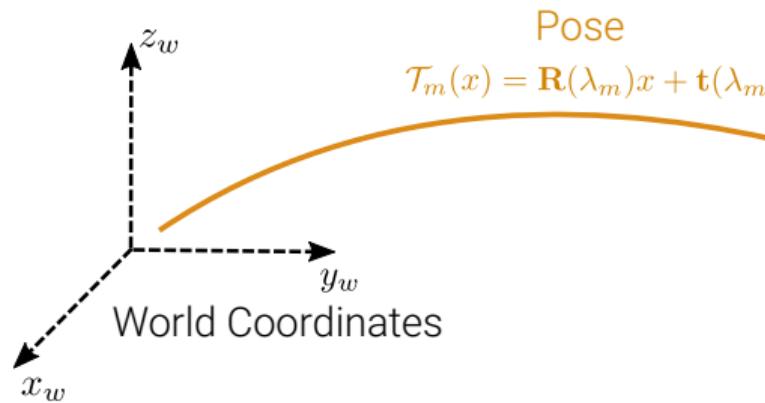


Pentland's Superquadrics Revisited



- ▶ 1 superquadric = 11 parameters \Rightarrow scene on the left stored in only 1000 bytes!
- ▶ However, early fitting-based approaches did not work robustly \Rightarrow learn from data

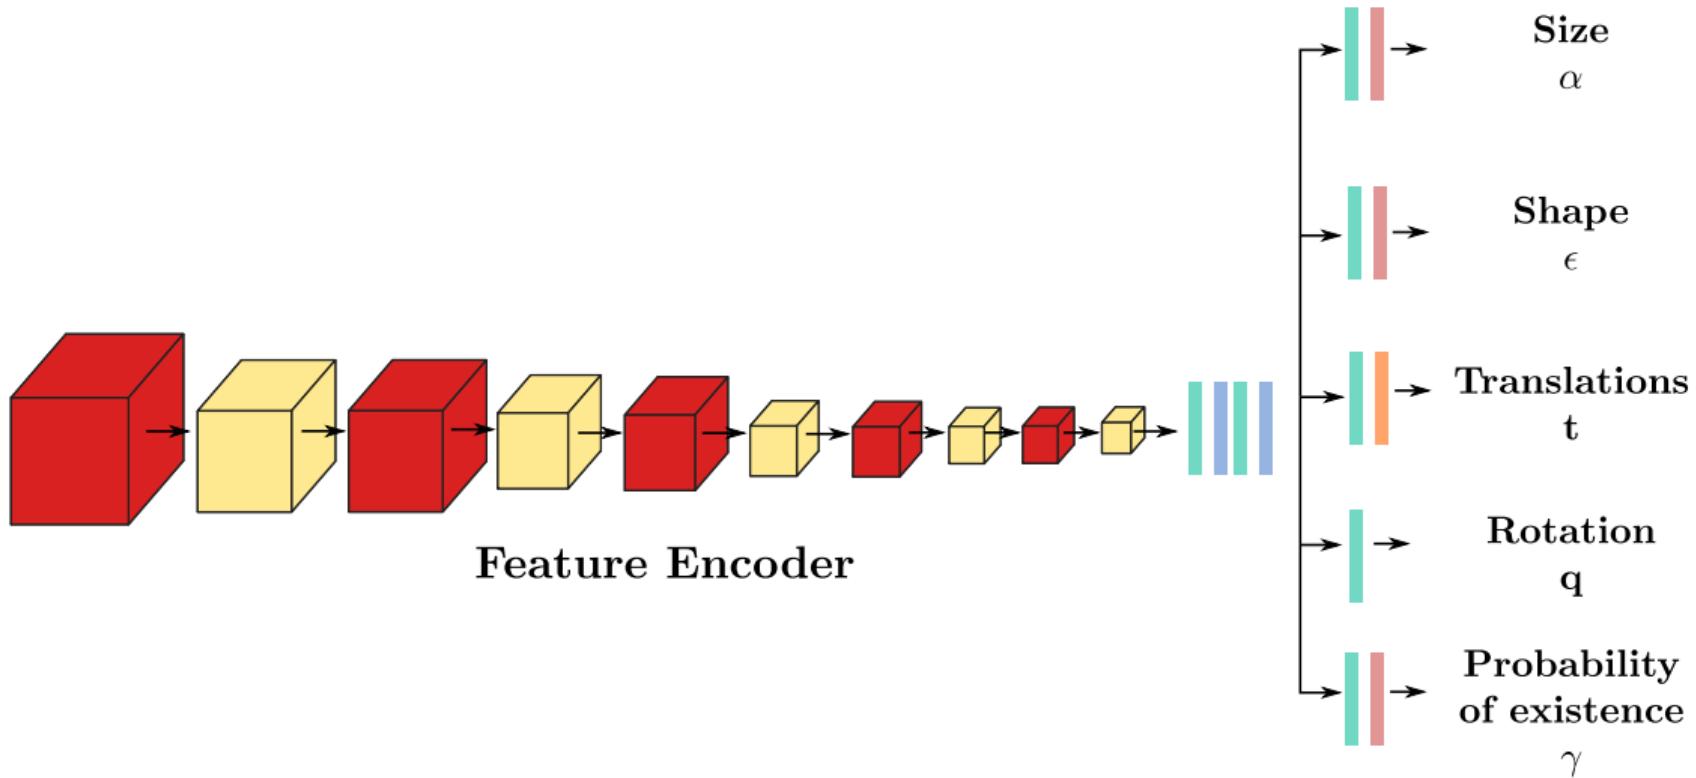
Learning 3D Shape Parsing



Predictions per primitive:

- ▶ 11 degrees of freedom: 6 pose (\mathbf{R} , \mathbf{t}) + 3 scale ($\boldsymbol{\alpha}$) + 2 shape ($\boldsymbol{\epsilon}$)
- ▶ Probability of existence: $\gamma \in [0, 1]$

Network Architecture



Loss Function

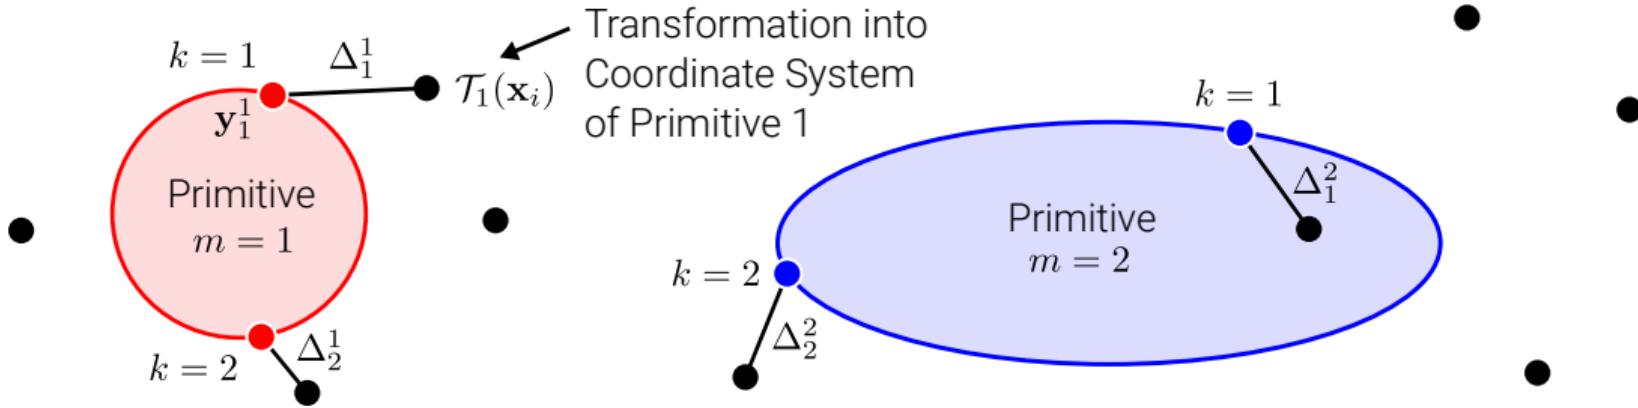
Overall Loss:

$$\mathcal{L}(\mathbf{P}, \mathbf{X}) = \mathcal{L}_{P \rightarrow X}(\mathbf{P}, \mathbf{X}) + \mathcal{L}_{X \rightarrow P}(\mathbf{X}, \mathbf{P}) + \mathcal{L}_\gamma(\mathbf{P})$$

Composed of:

- ▶ $\mathcal{L}_{P \rightarrow X}(\mathbf{P}, \mathbf{X})$: Primitive-to-Pointcloud Loss
- ▶ $\mathcal{L}_{X \rightarrow P}(\mathbf{X}, \mathbf{P})$: Pointcloud-to-Primitive Loss
- ▶ $\mathcal{L}_\gamma(\mathbf{P})$: Existence (at least one primitive) and Parsimony (sparsity) Loss

Primitive-to-Pointcloud Loss



$$\mathcal{L}_{P \rightarrow X}^m(\mathbf{P}, \mathbf{X}) = \frac{1}{K} \sum_{k=1}^K \Delta_k^m$$

$$\Delta_k^m = \min_{i=1,..,N} \|\mathcal{T}_m(\mathbf{x}_i) - \mathbf{y}_k^m\|_2$$

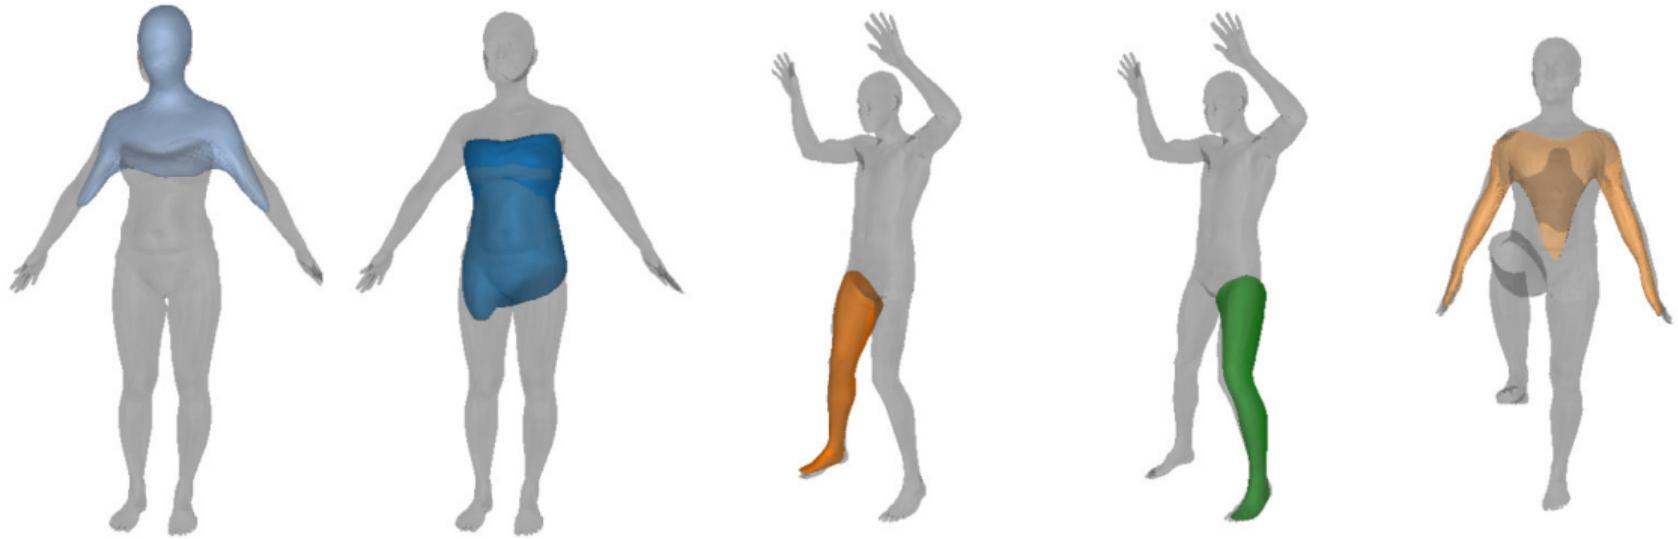
$$\begin{aligned} \mathcal{L}_{P \rightarrow X}(\mathbf{P}, \mathbf{X}) &= \mathbb{E}_{p(\mathbf{z})} \left[\sum_{m|z_m=1} \mathcal{L}_{P \rightarrow X}^m(\mathbf{P}, \mathbf{X}) \right] \\ &= \sum_{m=1}^M \gamma_m \mathcal{L}_{P \rightarrow X}^m(\mathbf{P}, \mathbf{X}) \end{aligned}$$

Results



<https://autonomousvision.github.io/superquadrics-revisited/>

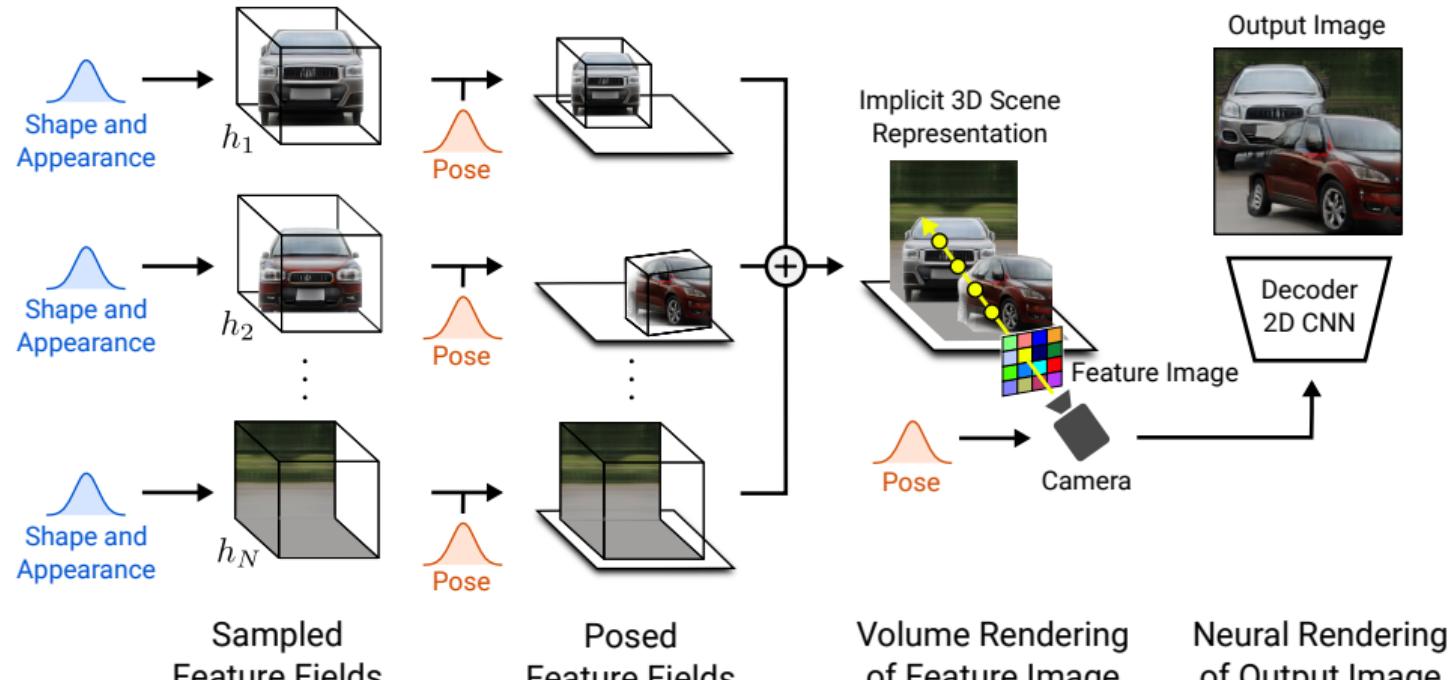
Neural Parts



- ▶ More **expressive part shapes** learned using invertible neural networks
- ▶ <https://autonomousvision.github.io/neural-parts/>

Compositional Feature Fields

GIRAFFE: Compositional Generative Neural Feature Fields



<https://autonomousvision.github.io/giraffe/>

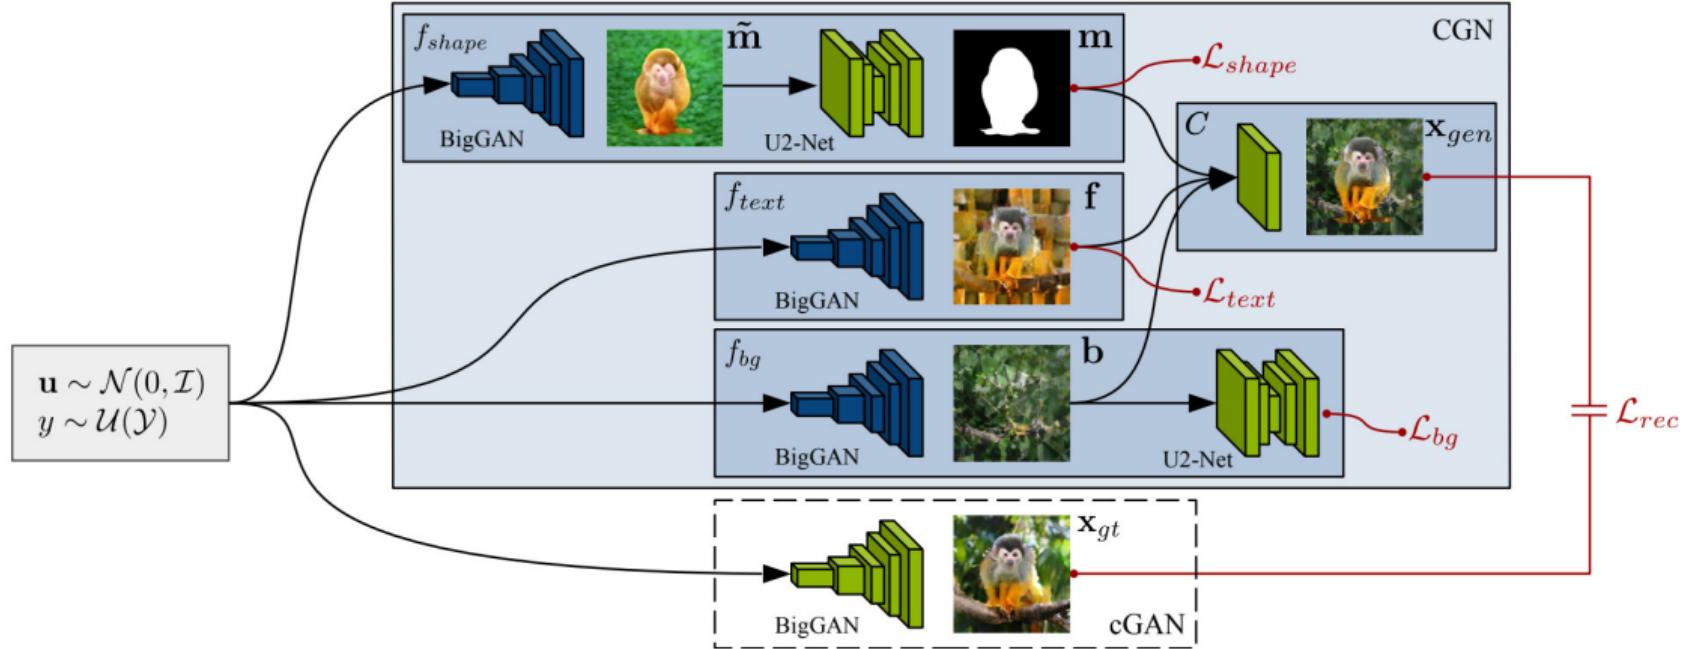
Causal Reasoning

Counterfactual Generative Networks



- ▶ Neural networks learn **spurious correlations** and often take **shortcuts**
- ▶ Can we learn to decompose image generation into **causal mechanisms?**
- ▶ This would allow to ask **counterfactual questions**, e.g.,
“How would this image look like with a different background?”
- ▶ Training on counterfactual data increases classifier **robustness** to OOD data

Counterfactual Generative Networks



Counterfactual Generative Networks

Shape	red wine	cottontail rabbit	pirate ship	triumphal arch	mushroom	hyena dog
Texture	carbonara	head cabbage	banana	Indian elephant	barrel	school bus
Background	baseball	valley	bittern (bird)	viaduct	grey whale	snorkel

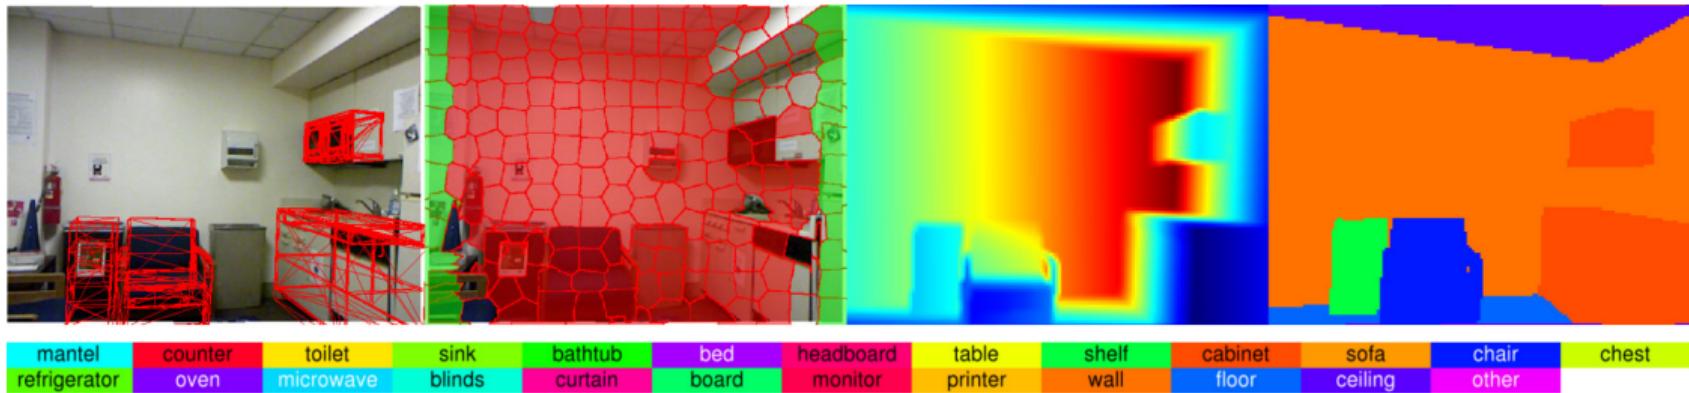


Figure 4: **ImageNet Counterfactuals.** The CGN successfully learns the disentangled shape, texture, and background mechanisms, and enables the generation of numerous permutations thereof.

<https://autonomousvision.github.io/cgn/>

Holistic 3D Scene Understanding

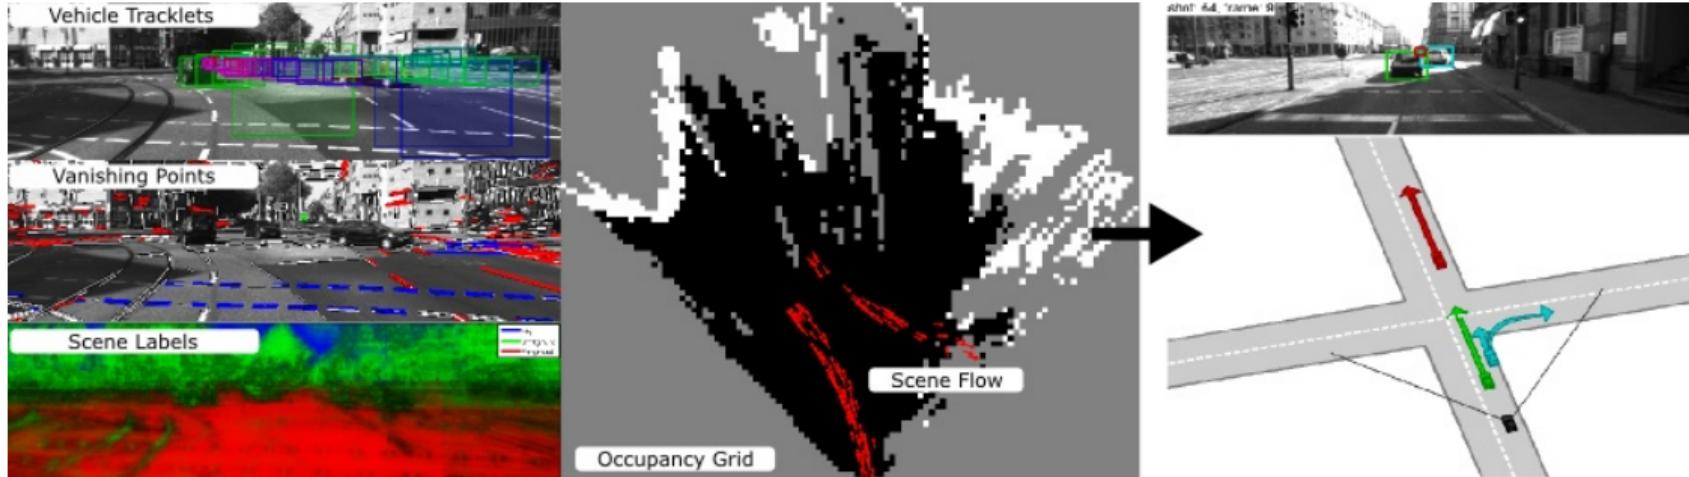
Indoor Scene Understanding



http://www.cvlabs.net/projects/indoor_scenes/

- Infer **3D objects and the layout** of indoor scenes from a single RGB-D image

Urban Scene Understanding



<http://www.cvlibs.net/projects/intersection/>

- ▶ Multi-object **traffic scene understanding** from movable platforms
- ▶ Reason about the 3D scene layout, the location of objects and traffic activities

12.3

Human Body Models

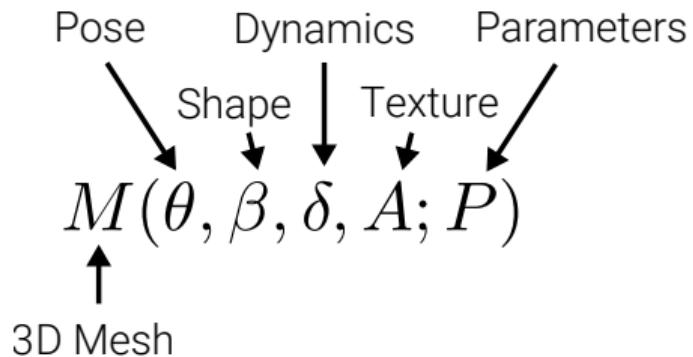
Humans Interact through their Bodies



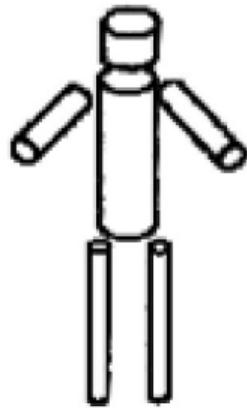
Generative Human Body Model



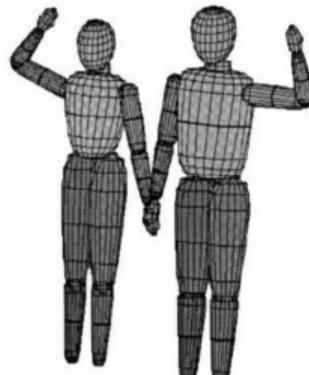
- ▶ Look like **real people**
- ▶ Deform like real people
- ▶ Small number of parameters
- ▶ **Easy to fit** to data
- ▶ **Easy to animate**



Early Body Models

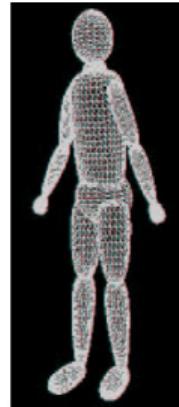


Marr and Nishihara '78
Nevatia and Binford '73

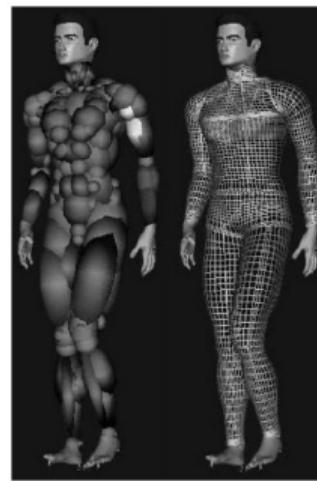


[Sminchisescu
and Triggs '03]

[Gavrilla, '96]



[Terzopoulos
and Metaxas '93]



[Plankers and
Fua '01]

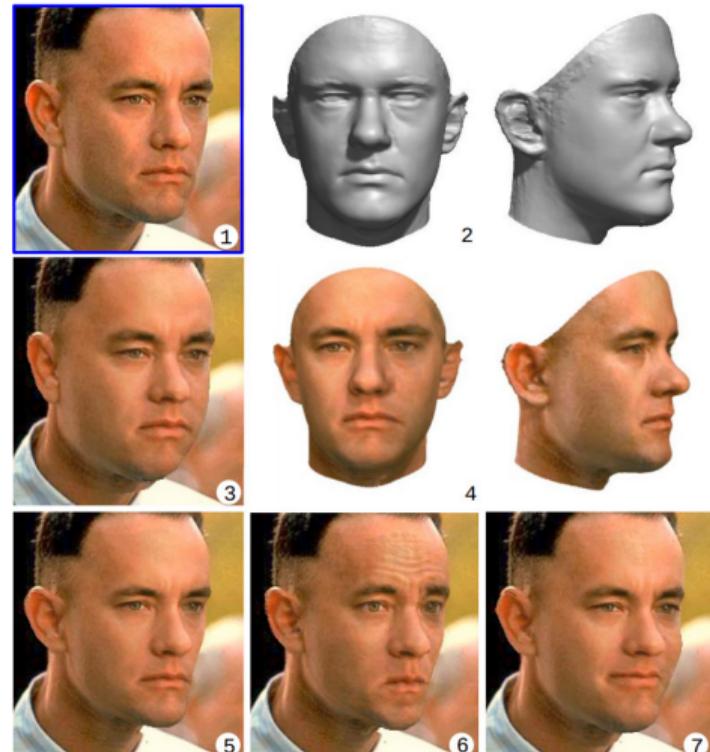
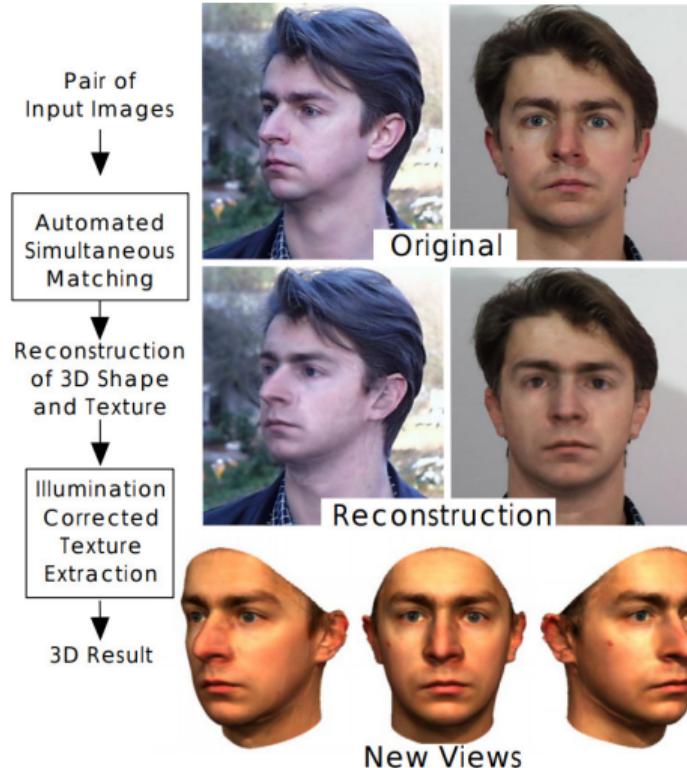


[Kakadiaris and Metaxas '00]



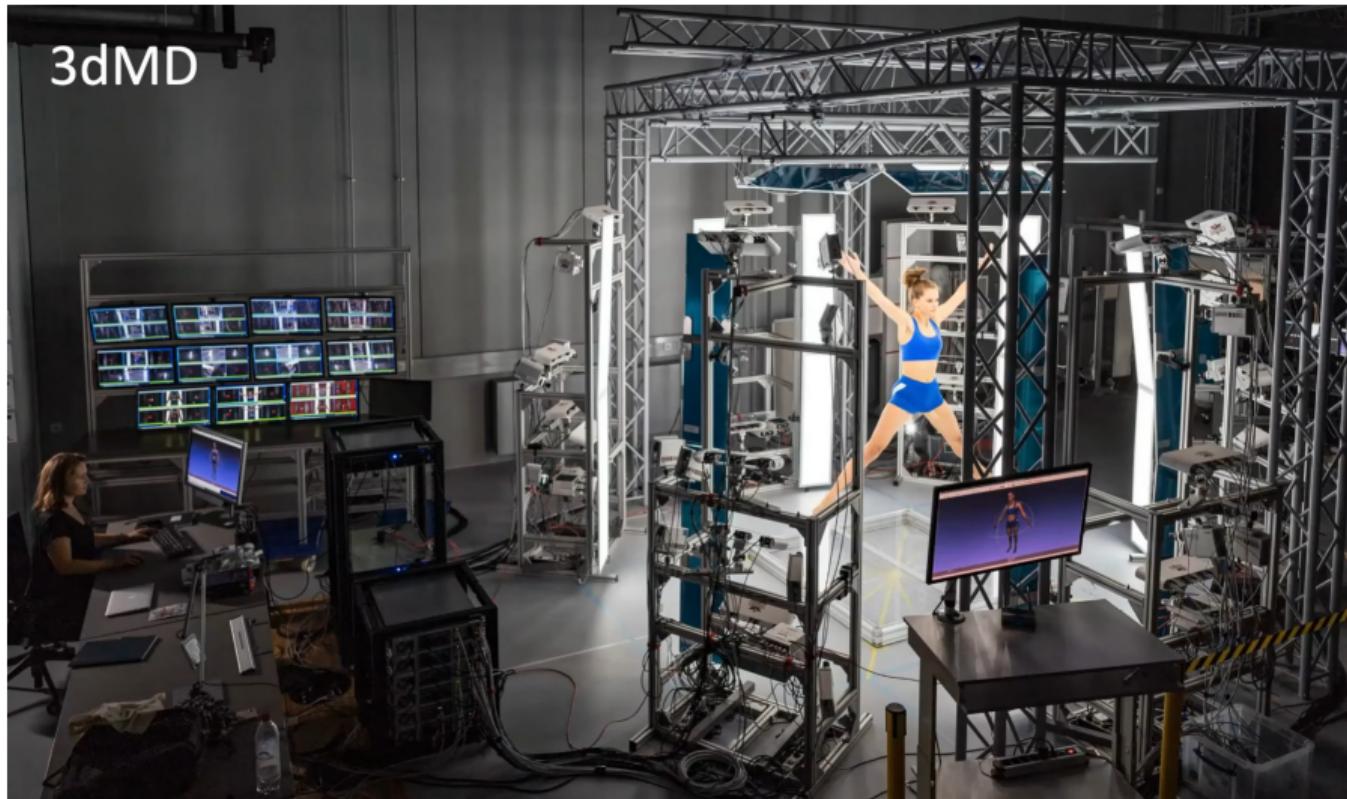
Nevatia & Binford '73

Inspiration: Morphable 3D Face Model

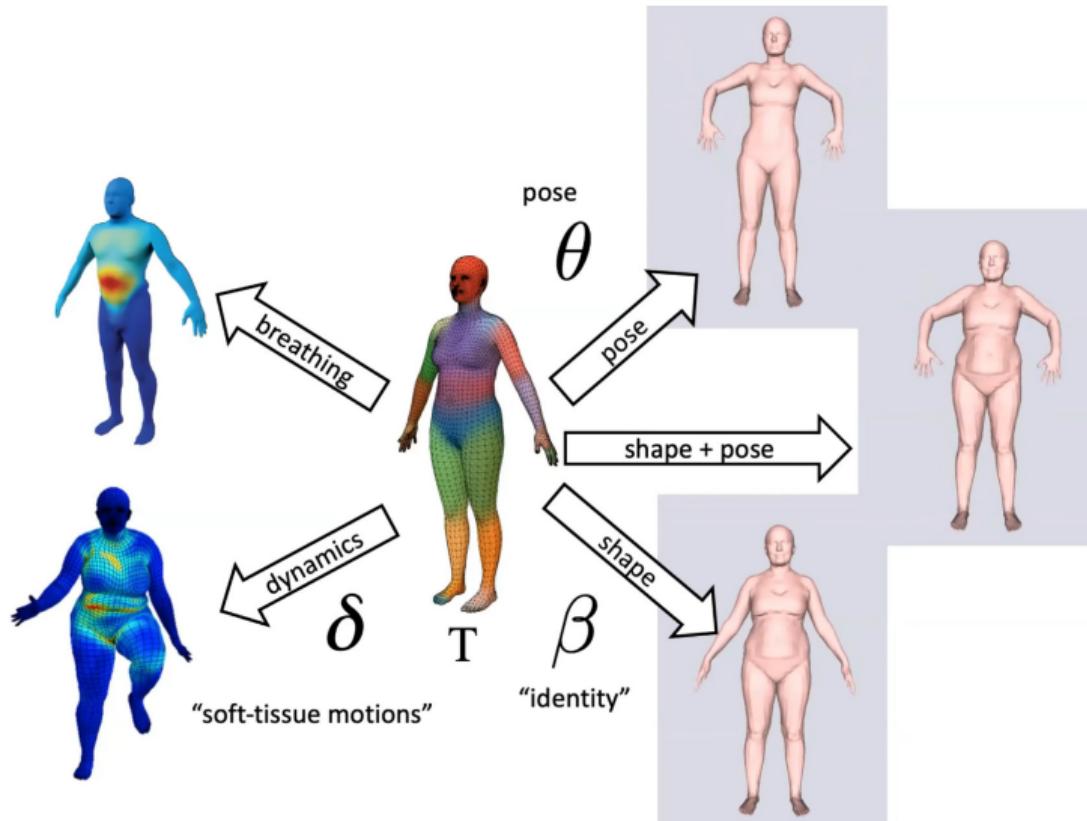


SMPL: A Skinned Multi-Person Linear Model

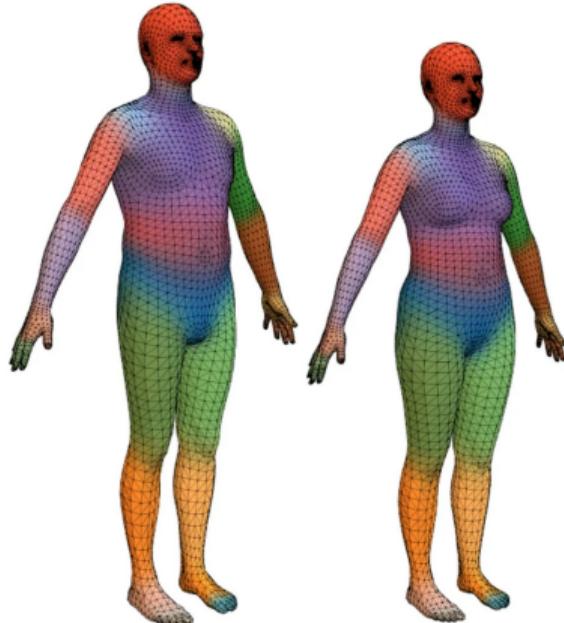
Input: Raw 3D Scans



Output: Factored Model

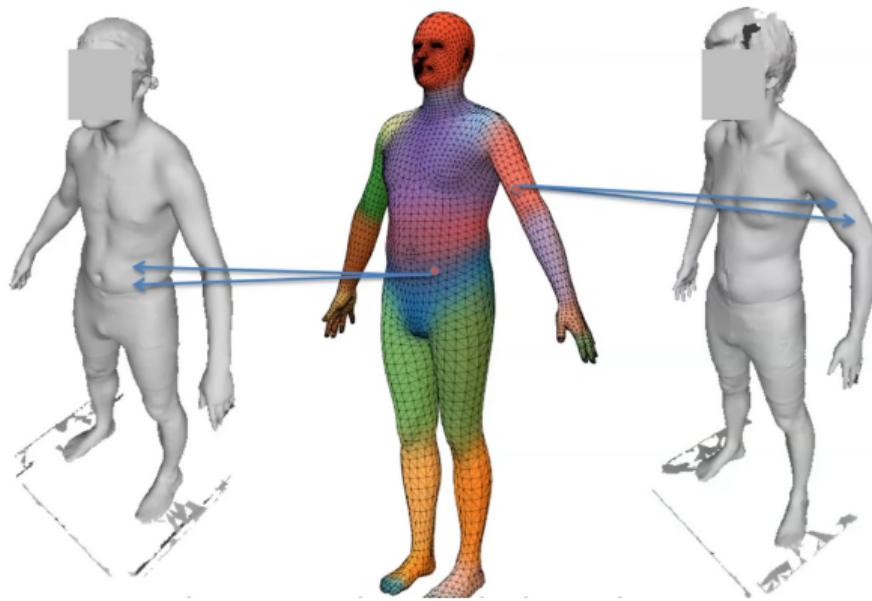


Representation: 3D Mesh



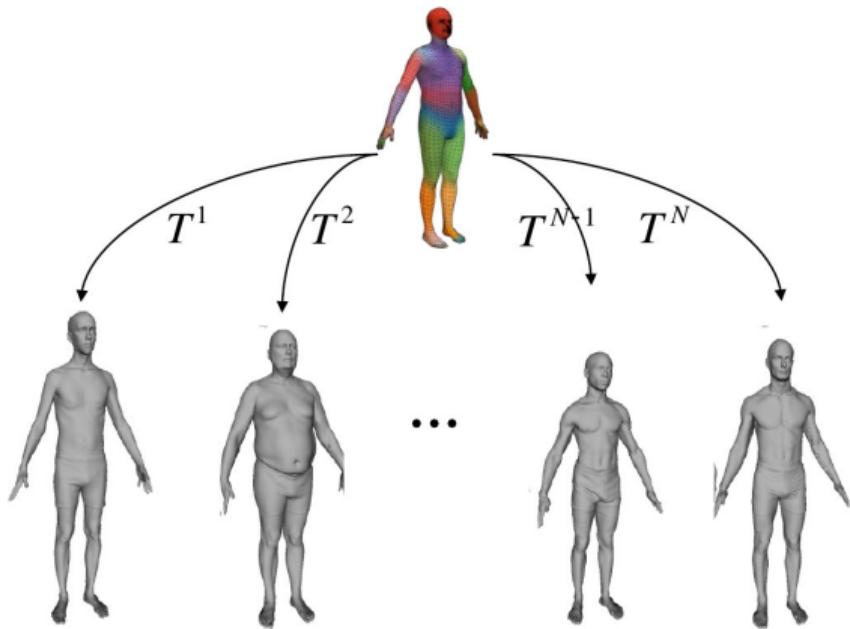
- ▶ SMPL uses a **3D mesh** as representation (vertices and faces)
- ▶ Base mesh designed by artist
- ▶ 7000 vertices \Rightarrow **21000 numbers**
- ▶ Most settings of these numbers do not correspond to people or shapes

Data Preprocessing: Mesh Registration



- ▶ Raw scans are unordered and incomplete 3D points
- ▶ To get correspondences, a template mesh must be **aligned** with each scan
- ▶ This is a hard problem
- ▶ **Chicken-and-egg problem**
(requires body model)
- ▶ Solution: solve model and registration **jointly**
- ▶ <http://faust.is.tue.mpg.de/>

Shape Deformation Subspace

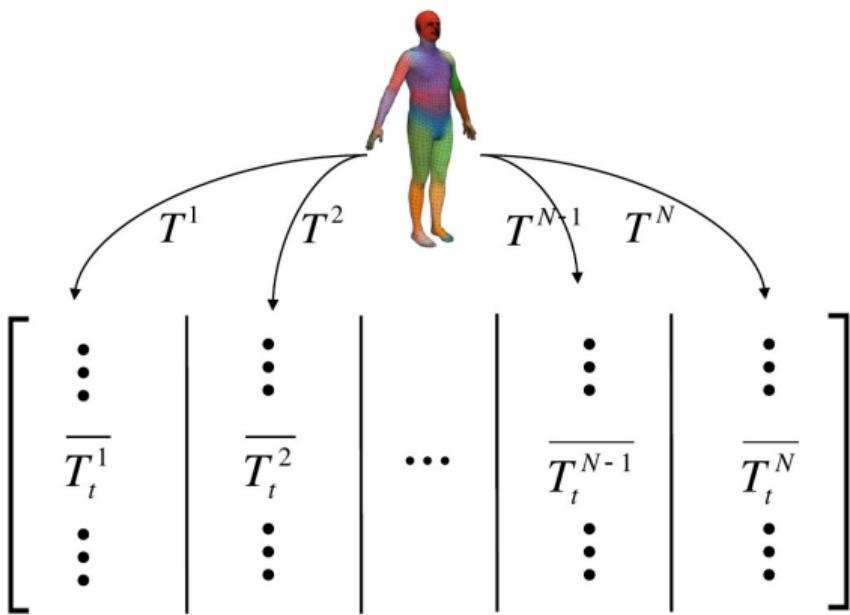


- ▶ Vectorize mesh vertices \mathbf{T}
- ▶ Subtract mean mesh (colored)
- ▶ Turn observations into matrix
- ▶ **Principal component analysis**
(PCA) yields a linear model:

$$\mathbf{T} = \mathbf{S}\boldsymbol{\beta} + \boldsymbol{\mu}$$

- ▶ Obtain low-dimensional subspace (10D-300D) that captures most of the variance in the 21000D space
- ▶ This works due to canonical pose

Shape Deformation Subspace

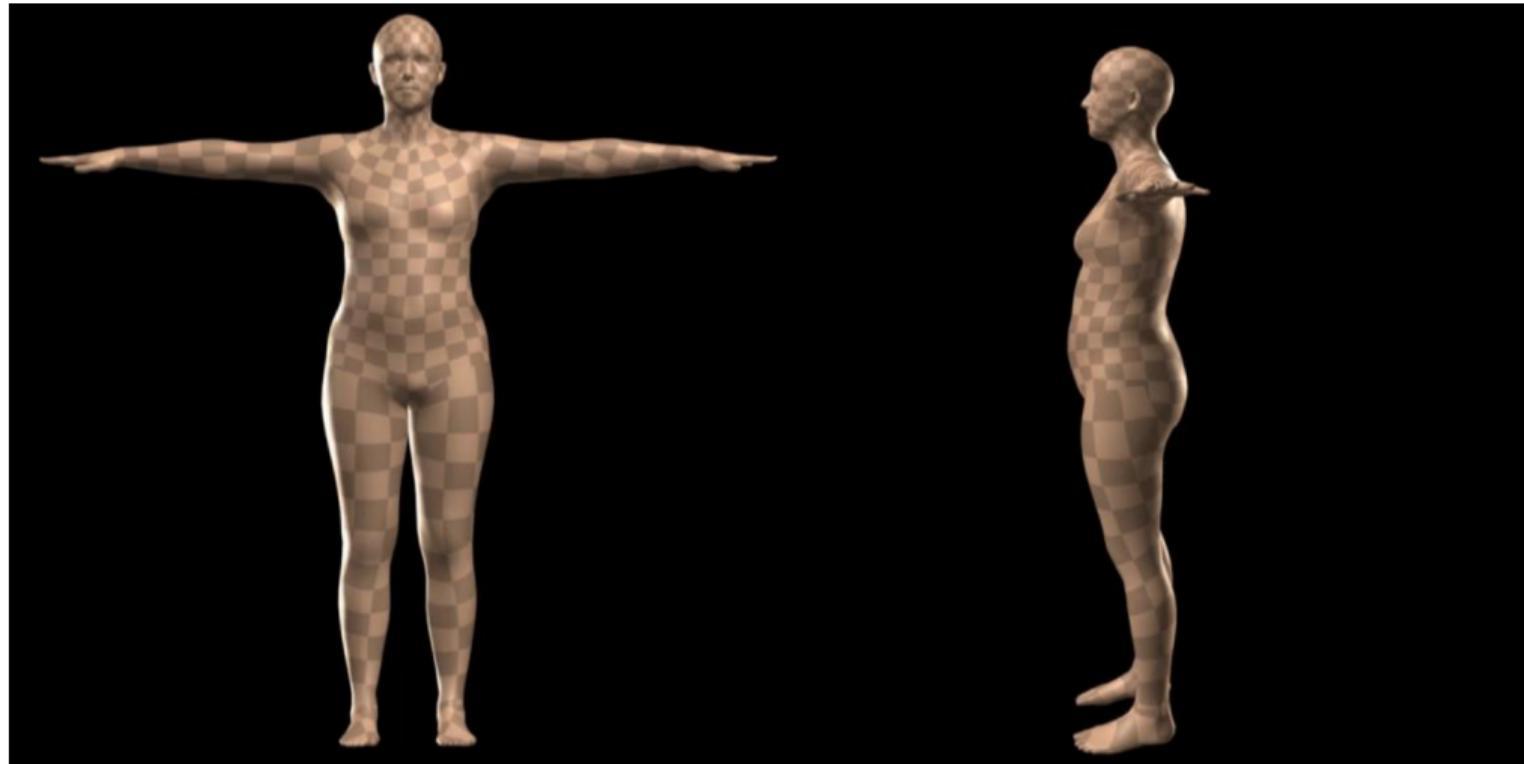


- ▶ Vectorize mesh vertices \mathbf{T}
- ▶ Subtract mean mesh (colored)
- ▶ Turn observations into matrix
- ▶ **Principal component analysis**
(PCA) yields a linear model:

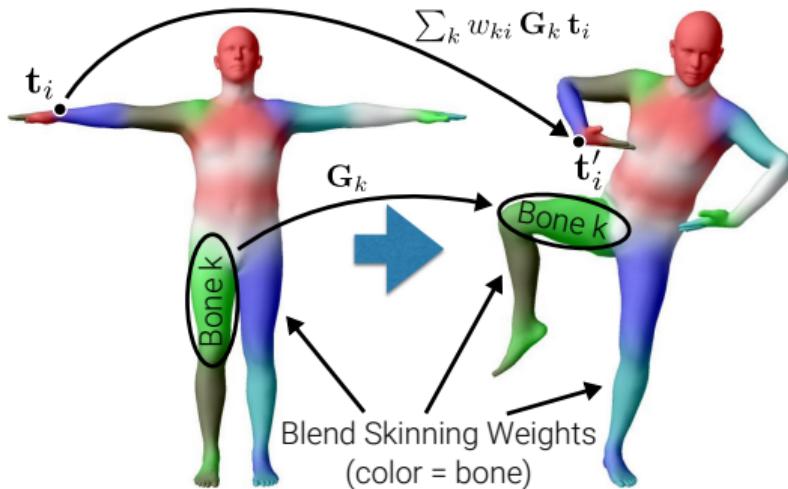
$$\mathbf{T} = \mathbf{S}\boldsymbol{\beta} + \boldsymbol{\mu}$$

- ▶ Obtain low-dimensional subspace (10D-300D) that captures most of the variance in the 21000D space
- ▶ This works due to canonical pose

Shape Deformation Subspace



Linear Blend Skinning



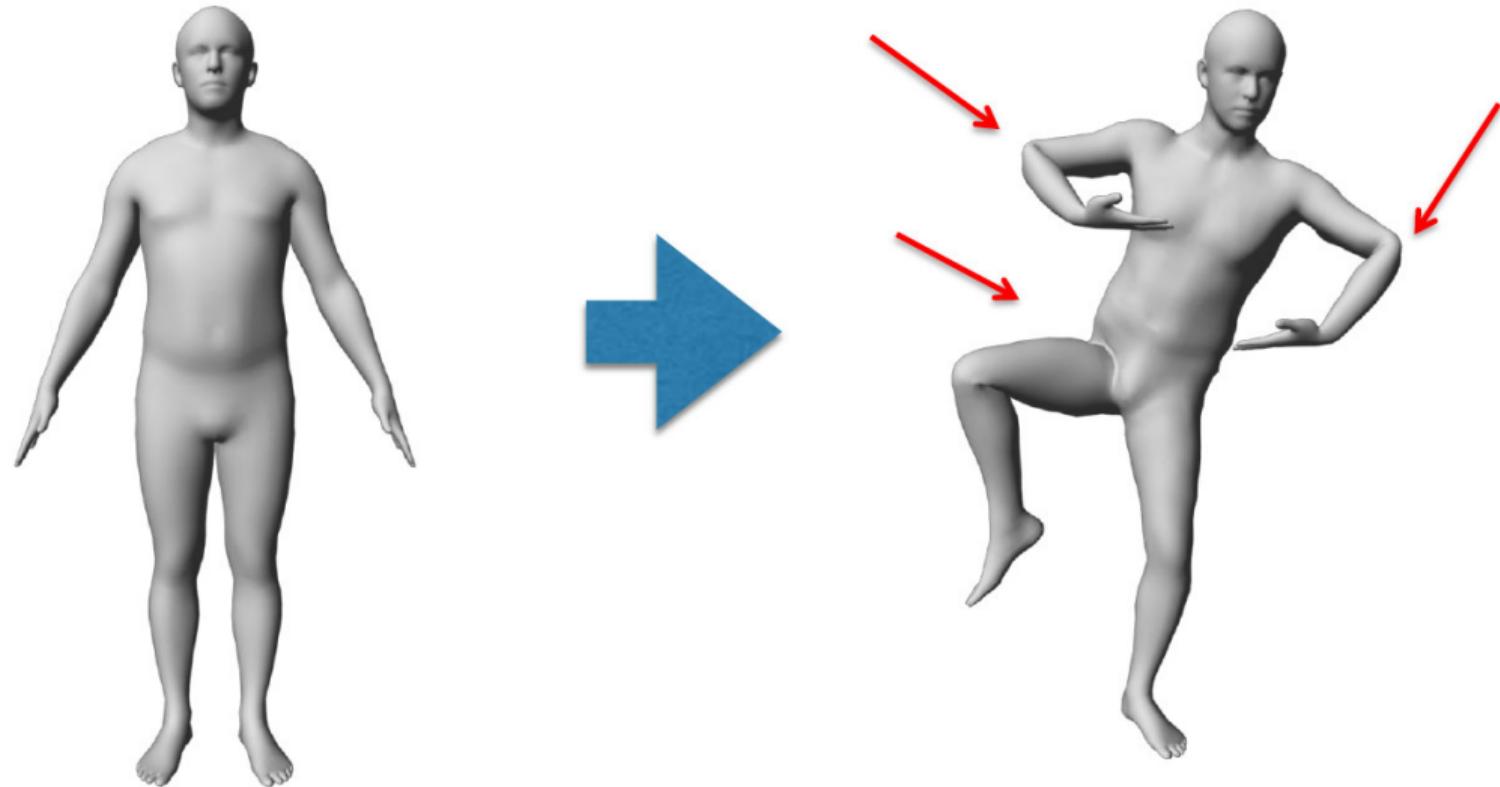
- Vertices are **linear combination** of transformed template vertices

Linear Blend Skinning (LBS):

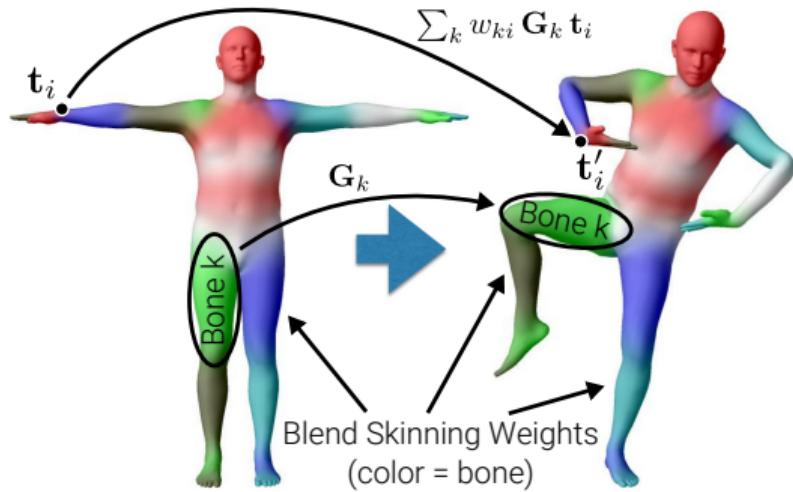
$$\mathbf{t}'_i = \sum_k w_{ki} \mathbf{G}_k(\theta, \mathbf{J}) \mathbf{t}_i$$

- \mathbf{t}, \mathbf{t}' : Rest/transformed vertices
- w_{ki} : Blend skinning weights
- \mathbf{G}_k : Rigid bone transformation
- θ : Pose
- \mathbf{J} : Joint locations
- Skinning weights created by artist

LBS Problems



SMPL



- Vertices are **linear combination** of transformed template vertices

LBS vs. SMPL:

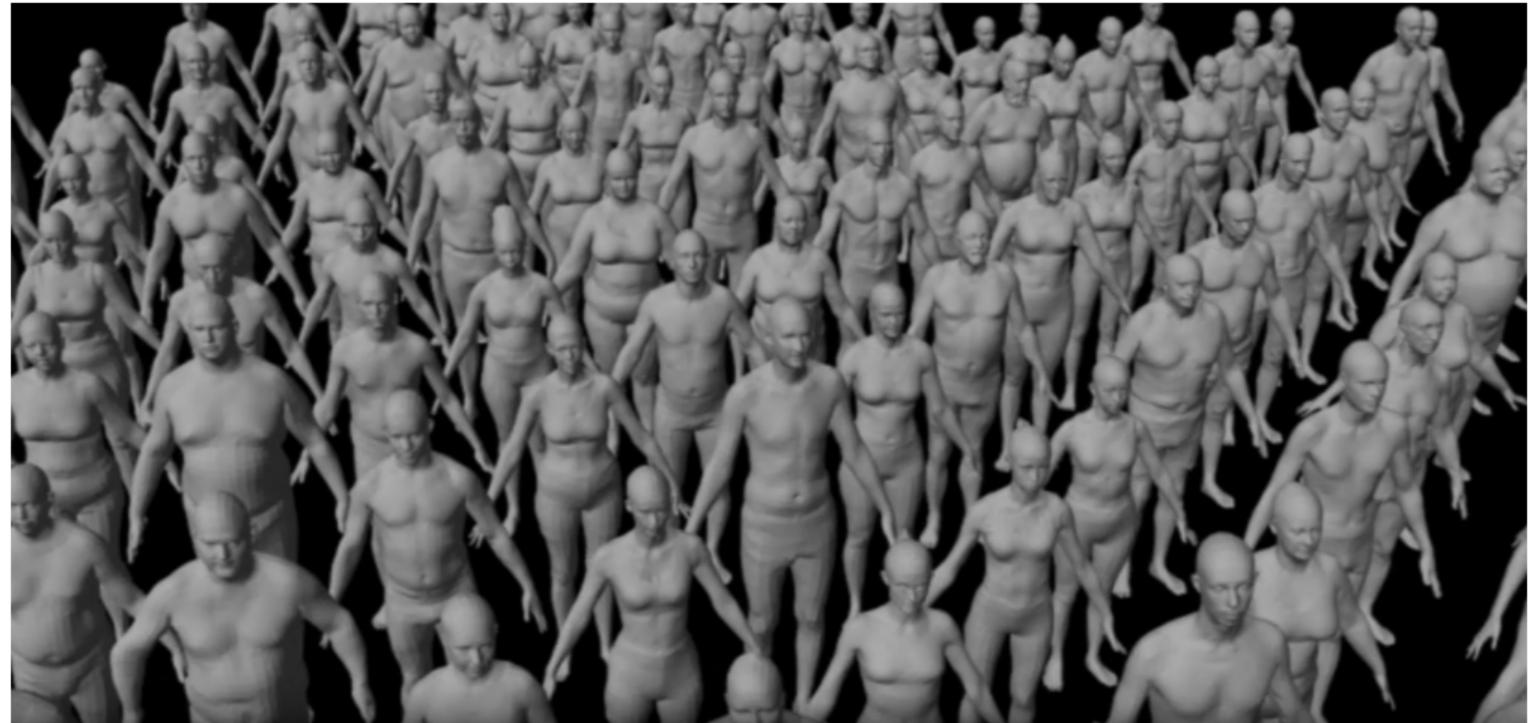
$$\mathbf{t}'_i = \sum_k w_{ki} \mathbf{G}_k(\boldsymbol{\theta}, \mathbf{J}) \mathbf{t}_i$$

$$\mathbf{t}'_i = \sum_k w_{ki} \mathbf{G}_k(\boldsymbol{\theta}, \mathbf{J}(\boldsymbol{\beta})) (\mathbf{t}_i + \mathbf{s}_i(\boldsymbol{\beta}) + \mathbf{p}_i(\boldsymbol{\theta}))$$

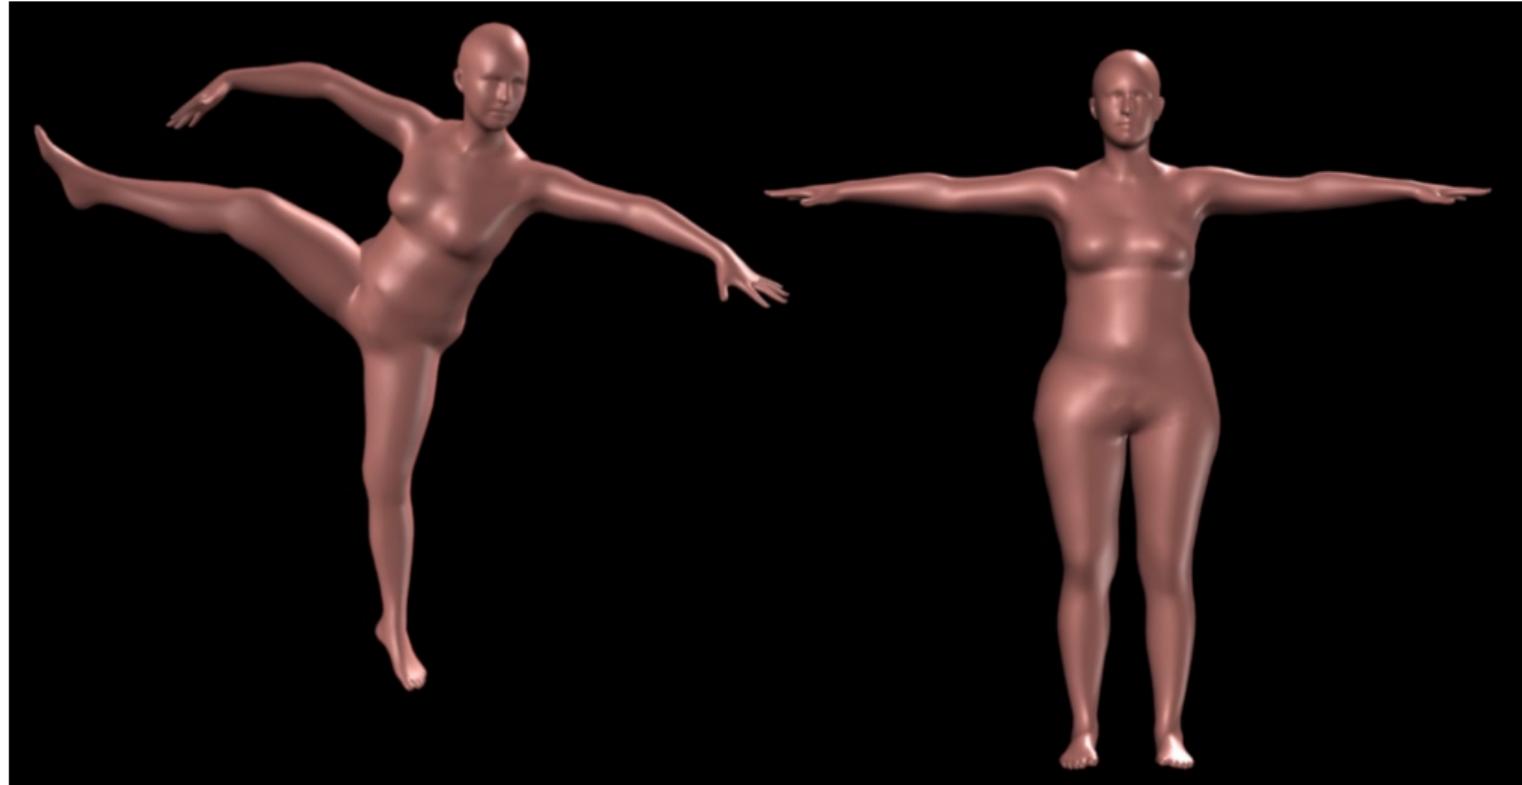
3 key differences:

- Joints $\mathbf{J}(\boldsymbol{\beta})$ depend on shape $\boldsymbol{\beta}$
- Shape correctives $\mathbf{s}(\boldsymbol{\beta}) = \mathbf{S}\boldsymbol{\beta}$
- Pose correctives $\mathbf{p}(\boldsymbol{\theta}) = \mathbf{P}\boldsymbol{\theta}$

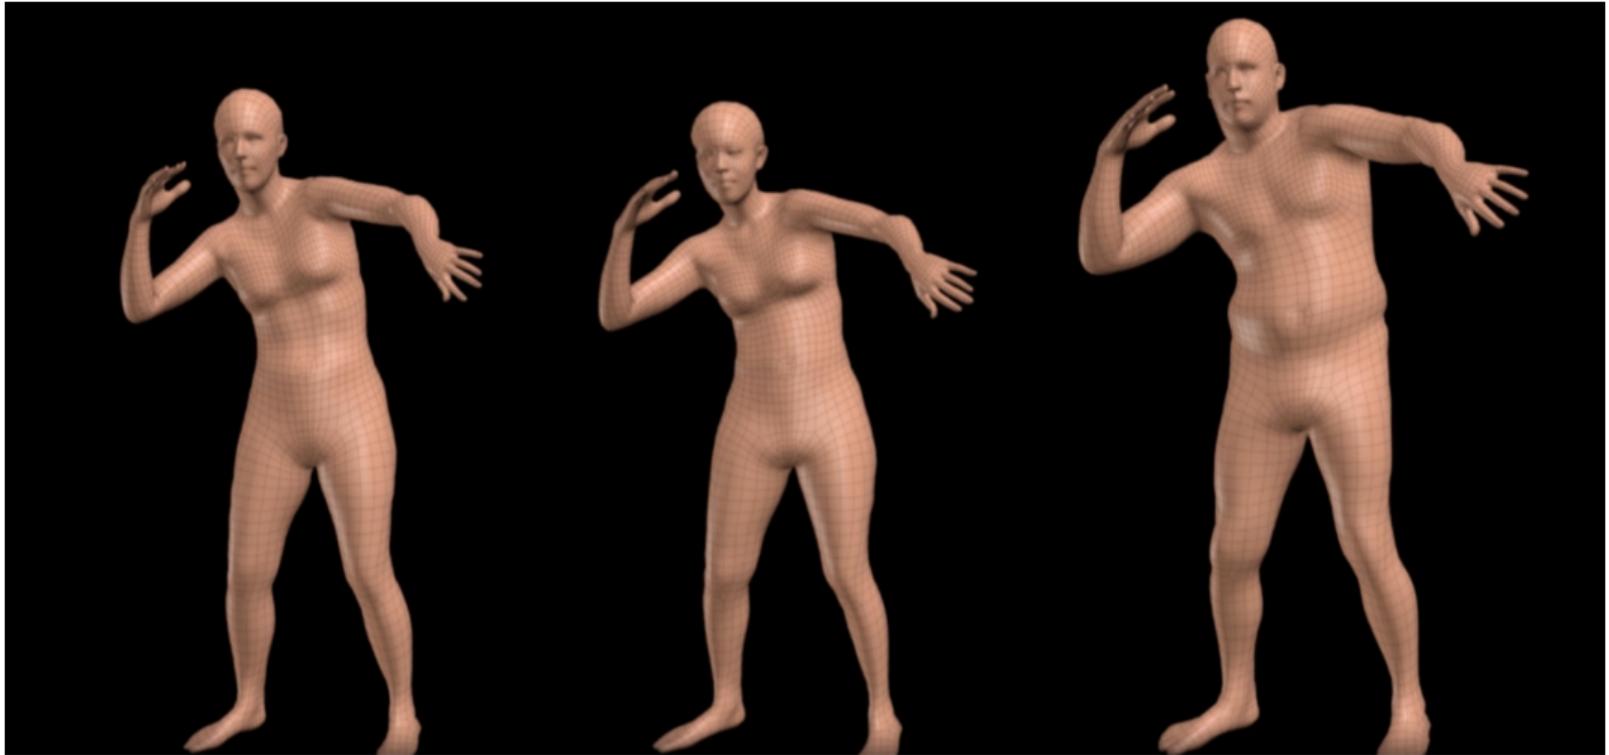
Learn Parameters from Registered Scans



Pose Blend Shapes



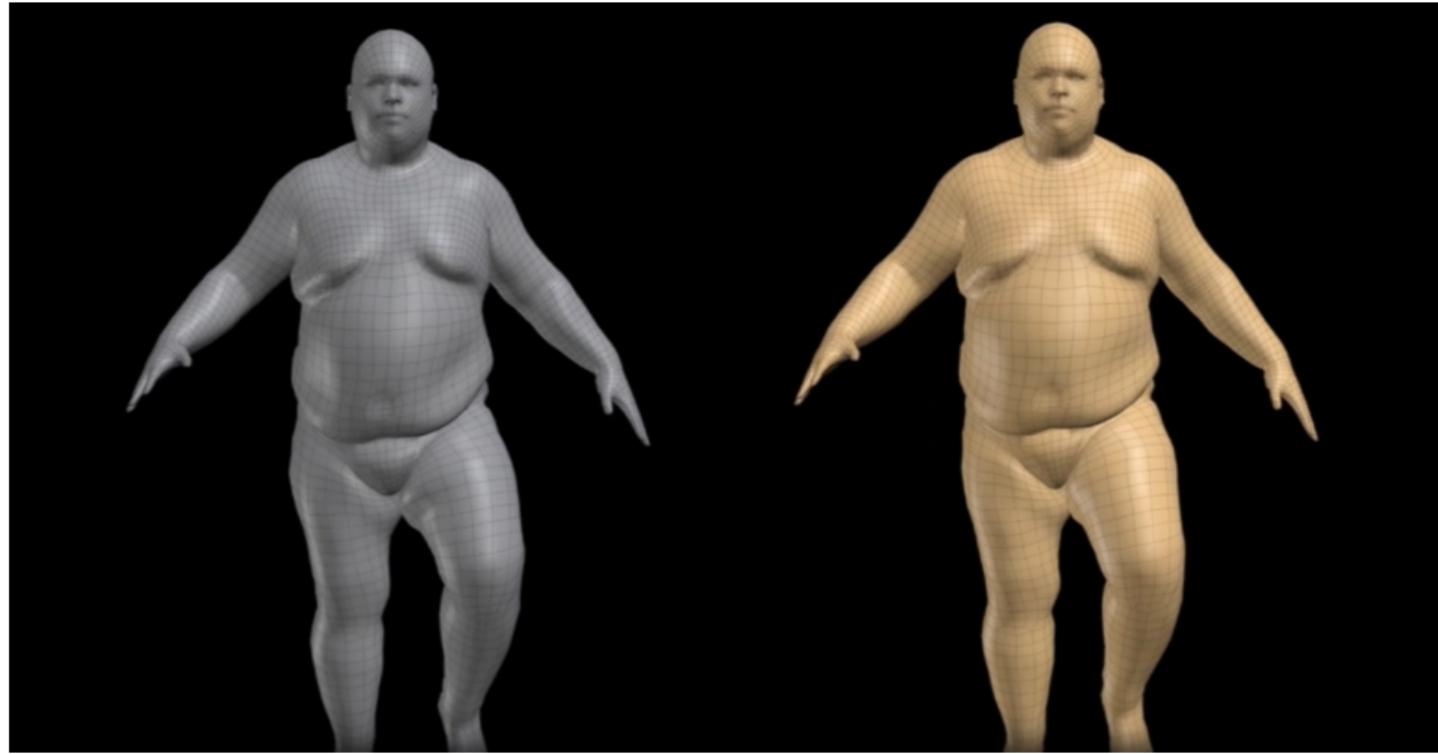
SMPL Results



Character Retargeting using SMPL



Dynamic SMPL (DMPL)



SMPLify: Estimation from a Single Image



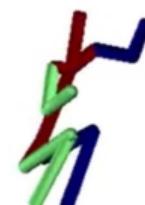
LSP



Akhter et al.[12]



Ramakrishna et al. [13]

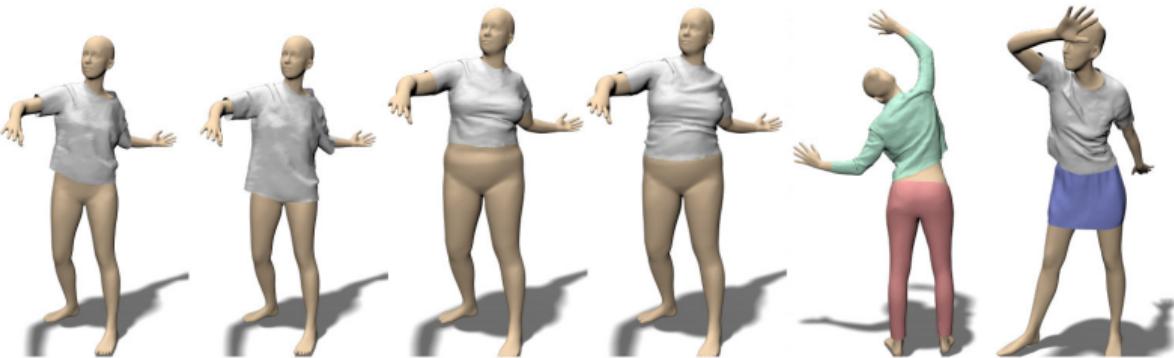


Zhou et al.[14]



SMPLify

Current Research Directions



- ▶ Clothing
- ▶ Garment
- ▶ Appearance
- ▶ Dynamics
- ▶ Few-shot
- ▶ From RGB
- ▶ Avatars
- ▶ Interaction
- ▶ Scenes

12.4

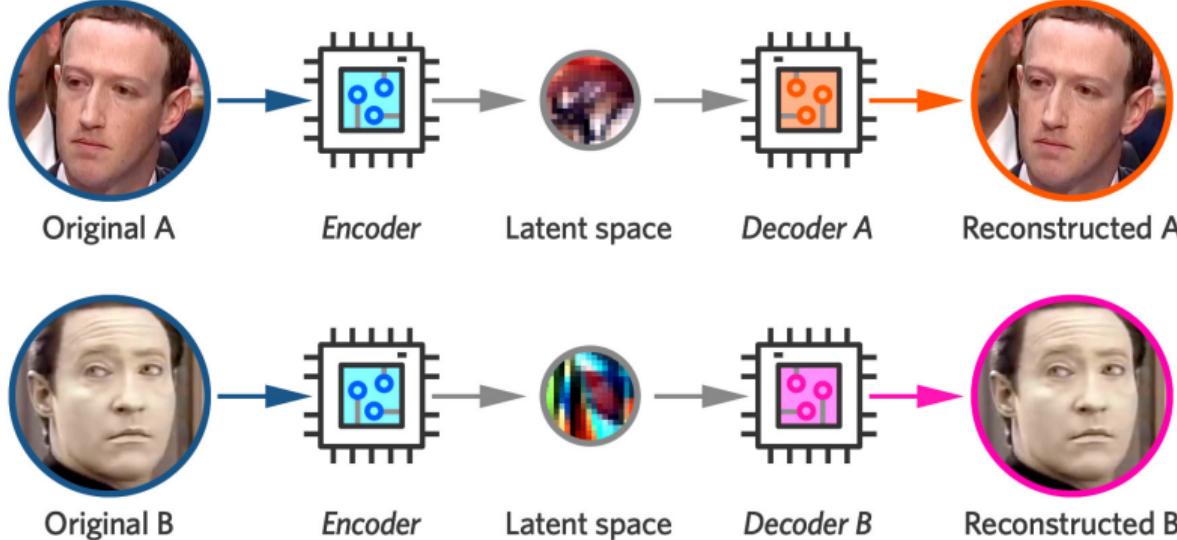
Deepfakes

What are DeepFakes?



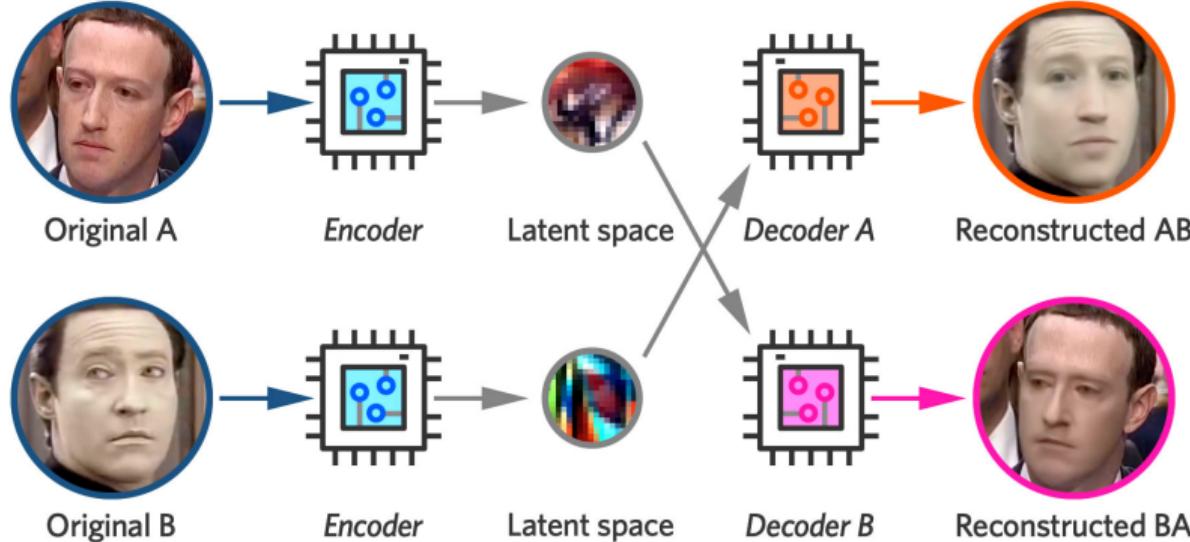
<https://github.com/deepfakes/faceswap>

FaceSwap



- ▶ Extract faces in two videos, train auto-encoder for each of the two persons
- ▶ Shared encoder weights to enforce same latent space (pose, expression, lighting)

FaceSwap



- ▶ Extract faces in two videos, train auto-encoder for each of the two persons
- ▶ Shared encoder weights to enforce same latent space (pose, expression, lighting)
- ▶ At inference time, swap decoders to decode latent code with new identity

FaceSwap



<https://arstechnica.com/science/2019/12/how-i-created-a-deepfake-of-mark-zuckerberg-and-star-treks-data/>

FaceSwap vs. Reenactment

Identity Swap



Facial Reenactment



- ▶ **FaceSwap:** Take face of video and copy it onto another video. Output is a hybrid between background (including hair) and new foreground (face).
- ▶ **Reenactment:** Take expression from one video and animate another person. Output is not a hybrid but manipulates expression of a person, enables animation.

Applications in Graphics and the Movie Industry

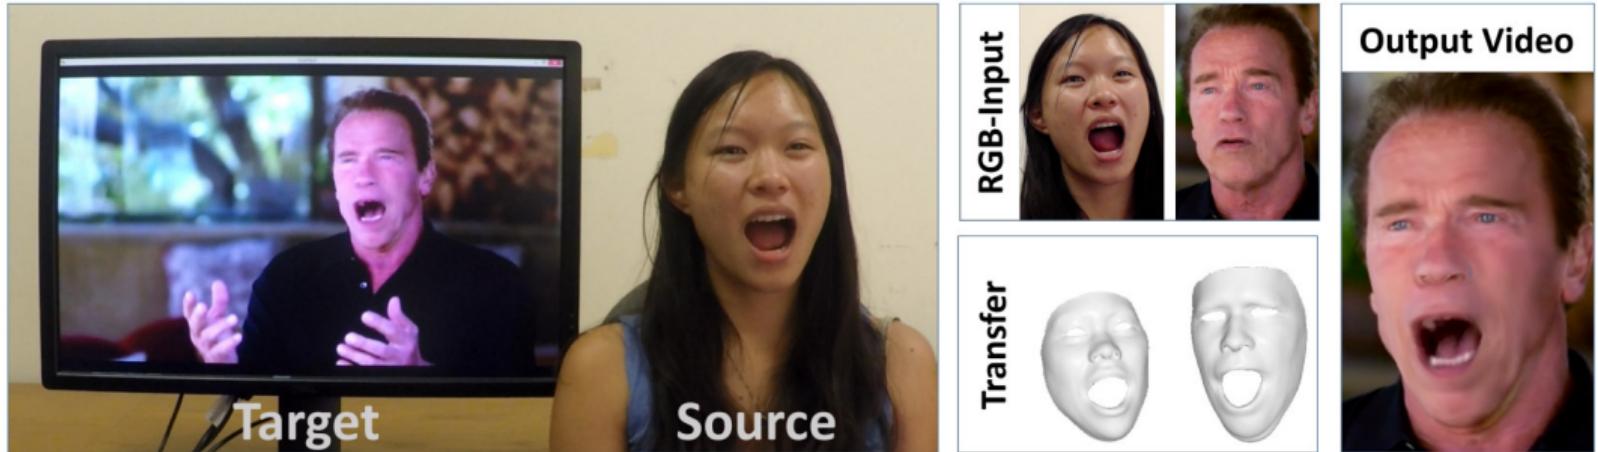


3D Model + Textures + Shading -> Synthetic Image



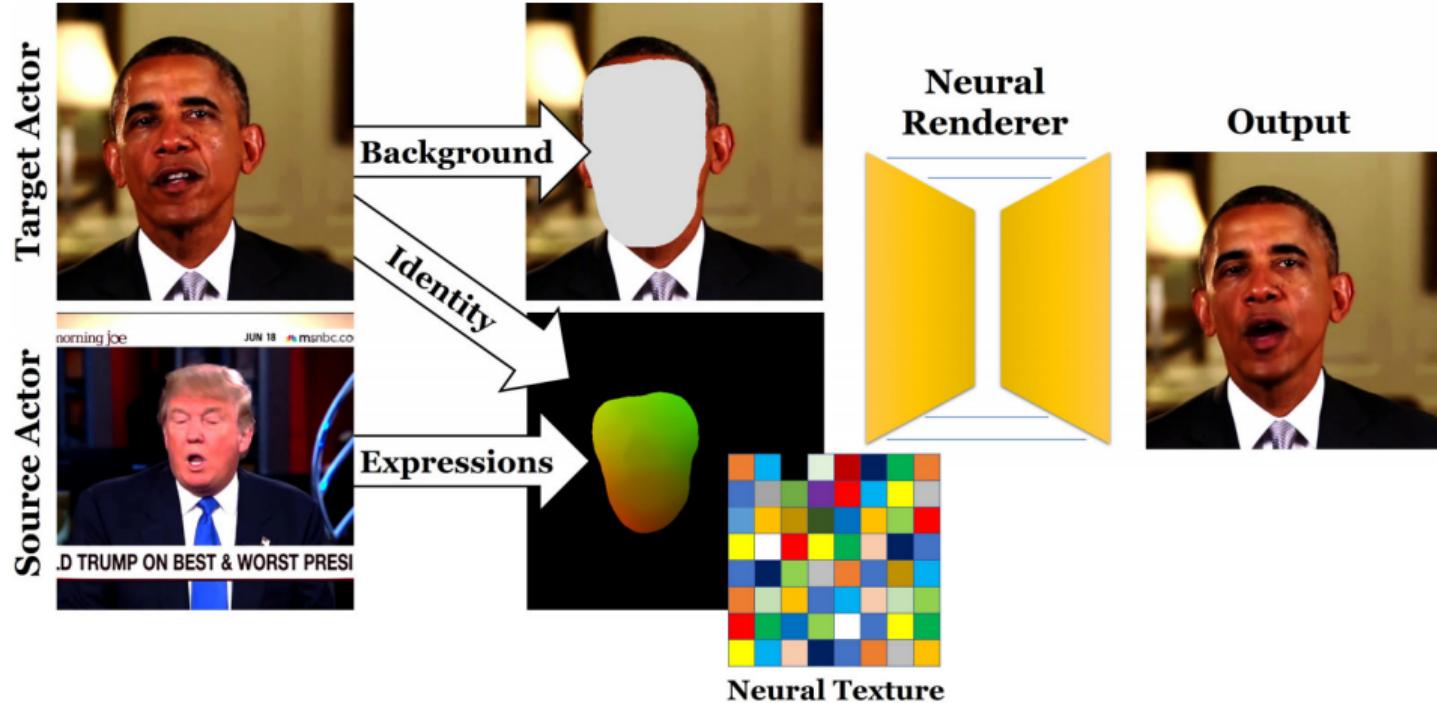
- ▶ Reenactment techniques have been used extensively in the **movie industry**
- ▶ But graphics methods require careful **3D model** creation and **manual editing**
- ▶ Learning has the potential to **automate this process** and enable richer editing

Face2Face



- ▶ **Real-time** facial reenactment, source video captured with commodity **webcam**
- ▶ **Face model** parameters recovered by non-rigid dense alignment wrt. image
- ▶ Track facial expressions of both source and target video for **appearance transfer**

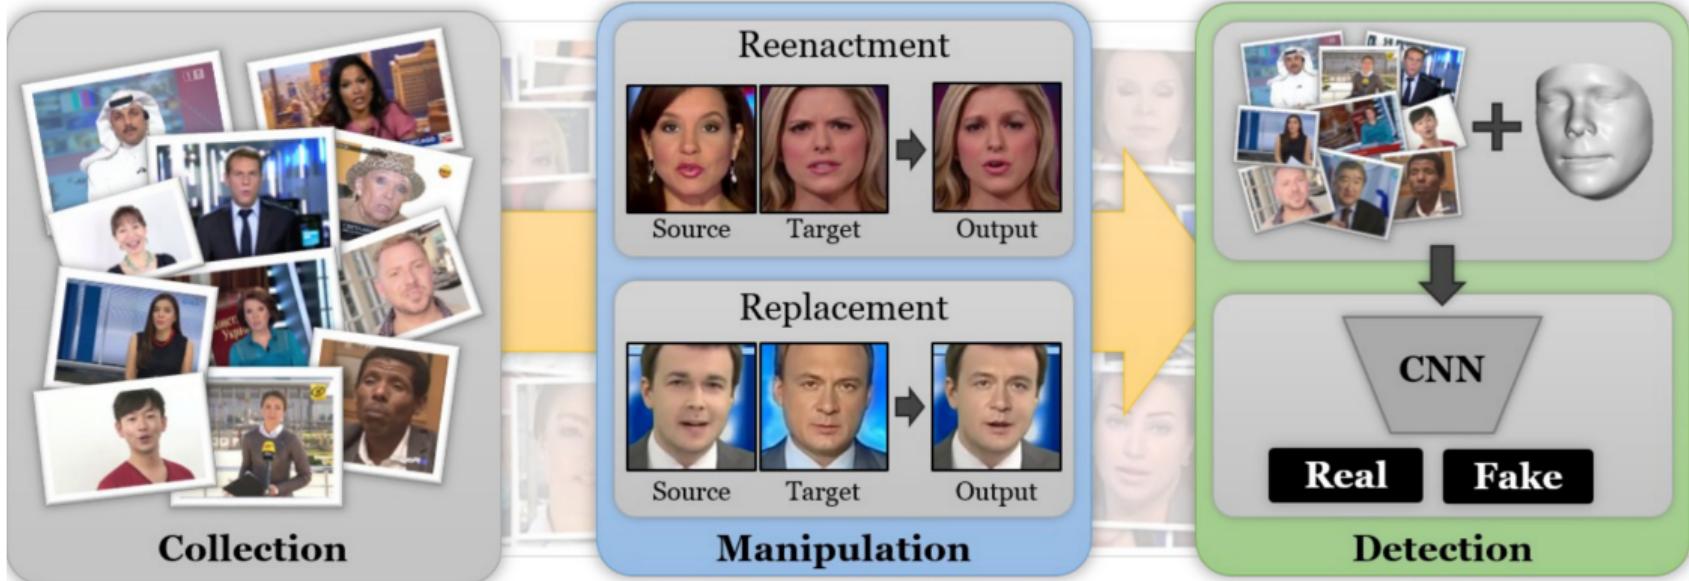
Deferred Neural Rendering



- ▶ Learn object-specific **neural textures** that can be “rendered” with a neural network

Deep Fake Detection

FaceForensics

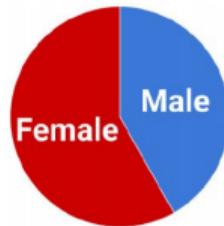


- Goal: Learn binary classifier to detect manipulated images

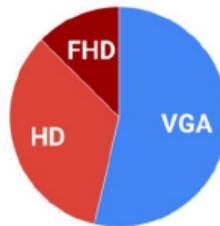
FaceForensics Dataset

Source: 1,000 Videos (510,529 frames)

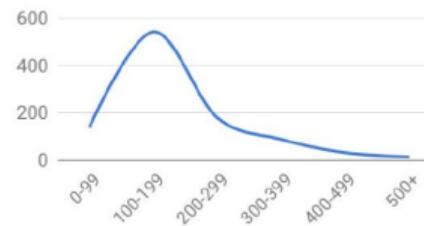
Methods	Train	Validation	Test
Pristine	366,847	68,511	73,770
DeepFakes	366,835	68,506	73,768
Face2Face	366,843	68,511	73,770
FaceSwap	291,434	54,618	59,640
NeuralTextures	291,834	54,630	59,672



(a) Gender



(b) Resolution



(c) Pixel Coverage of Faces

- Publicly available!
- Over 2 million manipulated frames
- Three compression levels for each manipulated frame
- Over 1000 research groups

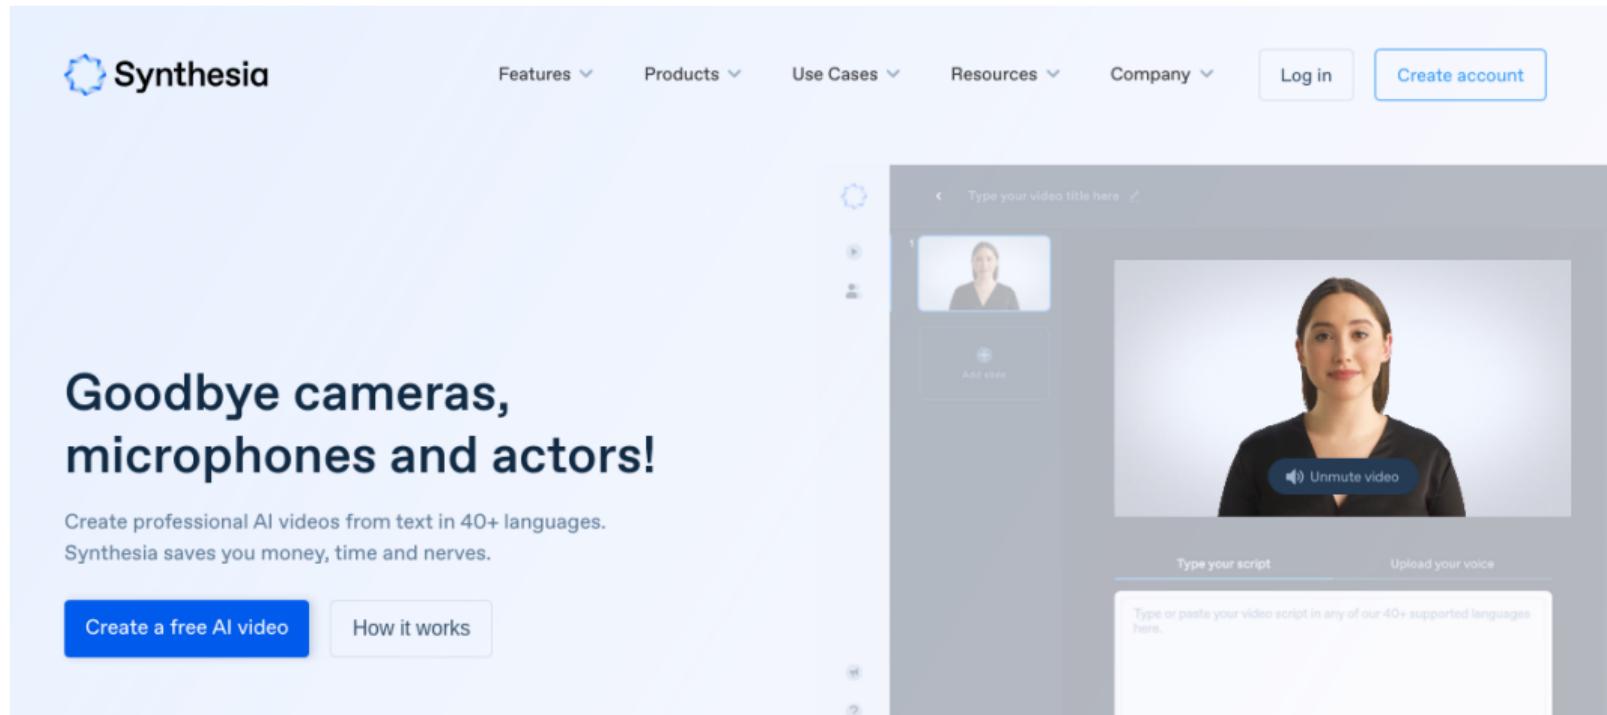
ForensicTransfer

XceptionNet	Test on Face2Face	Test on FaceSwap
Trained on Face2Face	98.13%	50.20%
Trained on FaceSwap	52.73%	98.30%

Conclusions:

- ▶ DeepFake detection works well on in-domain data (i.e., knowing the forgery method)
- ▶ However, it doesn't generalize
- ▶ Requires self-supervised, transfer and few-shot learning techniques for OOD generalization

Synthesia

The image shows the Synthesia website. At the top, there is a navigation bar with the Synthesia logo, a search bar, and links for Features, Products, Use Cases, Resources, Company, Log in, and Create account. Below the navigation, a large banner features the text "Goodbye cameras, microphones and actors!" in bold, dark blue font. Underneath this, two smaller text blocks read "Create professional AI videos from text in 40+ languages." and "Synthesia saves you money, time and nerves." To the right of the banner is a screenshot of the AI video creation interface. It shows a preview window with a woman's face, a script input area with placeholder text, and a "Type your voice" button. A "Unmute video" button is also visible.

Features ▼ Products ▼ Use Cases ▼ Resources ▼ Company ▼

Log in Create account

Goodbye cameras, microphones and actors!

Create professional AI videos from text in 40+ languages.
Synthesia saves you money, time and nerves.

Type your video title here

Add slide

Unmute video

Type your script

Type or paste your video script in any of our 40+ supported languages here.

Upload your voice

Create a free AI video How it works

<https://www.synthesia.io/>

Thank You!

All the best for the exam